

Verifying type- and scope-safe program transformations

Piotr Jander

MInf Project (Part 2) Report

Master of Informatics
School of Informatics
University of Edinburgh

2019

Abstract

There is an ongoing effort in the programming languages community to verify correctness of compilers. Type- and scope-safe representation is a commonly used encoding for intermediate languages, which however requires writing considerable meta-theoretical boilerplate for each IR. Standard techniques for showing correctness of compiler transformations are bisimulations and Kripke logical relations.

This project formalises a language with closures, implements a closure conversion algorithm, and mechanises the proof of its correctness using the the aforementioned techniques.

This project also suggest that a language with closures cannot benefit from state-of-the-art techniques for reducing meta-theoretical boilerplate with generic programming.

Acknowledgements

Acknowledgements go here.

Table of Contents

1	Introduction	5
2	Related literature	6
2.1	Verified compilation	6
2.2	Closure conversion	7
3	Background	8
3.1	Closure conversion	8
3.2	Compilation phases and intermediate representations	8
3.3	Type- and scope-safe representation of simply typed lambda calculus λ_{st}	9
3.4	Type- and scope-safe programs	11
3.5	ACMM's notion of a semantics	12
3.6	Small-step operational semantics	13
4	Formalising closure conversion	14
4.1	Closure language λ_{cl}	14
4.1.1	Terms	14
4.1.2	Renaming and substitution	15
4.1.3	Operational semantics	16
4.1.4	Conversion from λ_{st} to λ_{cl}	16
4.1.5	Minimising closure conversion	17
4.1.6	Agda implementation of minimising closure conversion	19
4.1.7	Fusion lemmas for the closure language λ_{cl}	21
5	Proving correctness of closure conversion with bisimulation	24
5.0.1	Compatibility relation	25
5.1	Bisimulation	26
5.2	Compatibility, values, renaming, and substitution	27
5.3	Compatibility relation is a bisimulation	29
5.4	Compatibility relation and closure conversion	29
6	Proof by logical relations	31
6.1	Alternative formalisation of the intermediate languages	31
6.2	Proving correctness of a translation with logical relations	33
7	Reflections and evaluation	37

8 Relationship to the UG4 project	41
8.1 Program transformation vs derivation	41
8.2 Program derivations in the UG4 project	42
8.3 Rewrite rules	42
8.4 Program derivations in compilers and their limitations	43
8.5 Implementation of rewriting in the UG4 project	43
8.6 Program derivation in Agda: sparse matrix-vector multiplication . . .	44
Bibliography	45
9 Technical appendix	47
9.1 Minimising closure conversion and the similarity relation	47
9.2 Compatibility relation is a bisimulation	49

Chapter 1

Introduction

TODO

Chapter 2

Related literature

2.1 Verified compilation

Closure conversion is just one possible verification phase, and its verification constitutes part of a wider effort to verify compilation end-to-end, which usually entails verifying operational correctness of all compilation phases.

As far as type safety is concerned, the reference is a paper by Morrisett et al., "From System F to Typed Assembly Language" [?]. It builds upon previous results in type safety of compilation phases (like the aforementioned [?]) and describes a typed RISC-like assembly (named TAL), which is the target of the final phases of compilation. As a whole, the paper proves type safety for a compilation pipeline from System F to TAL. It does not, however, prove end-to-end operational correctness.

An compiler which was verified for end-to-end operational correctness was described by Adam Chlipala in his paper "A Certified Type-Preserving Compiler from Lambda Calculus to Assembly Language". The source is a variant of the simply-typed lambda calculus (STLC). Compilation proceeds through six phases, eventually yielding idealised assembly code. The compiler is implemented in Coq, where terms and functions on terms are dependently typed, guaranteeing type preservation. This is also the approach taken in this project, except that we use Agda instead of Coq [?]. Operational correctness is proved by adopting denotational semantics, unlike in this project, which uses operational semantics. Due to unfamiliarity with operational semantics, we cannot comment on which approach is better (TODO or can we?).

Another example of a certified compiler is CompCert [?], which is the result of the first successful attempt to implement a certified compiler of a real-world (TODO wording) language. Even compared with the simply-typed lambda calculus (STLC), which was the source language in Chlipala's work [?], the C language is in some ways simpler, especially since it does not have first-class functions with free variables (TODO wording: scoping?). But, being a fully-fledged language, C presents enough challenges as the source language of a verified compiler.

2.2 Closure conversion

Closure conversion is a compilation phase where functions or lambda abstractions with free variables are transformed to /closures/. A closure consists of a body (code) and the /environment/, which is a record holding the values corresponding to the free variables in the body (code). Closure conversion transforms abstractions to closures, and replaces references to variables with lookups in the environment.

Closure conversion was necessarily used in every compiler for a language which supports functions with free variables (TODO wording: scope?). But the first work which provided a rigorous treatment of closure conversion was the paper "Typed Closure Conversion" by Minamide et al. [?]. It demonstrated type-preserving closure conversion, where closure environments have existential types (TODO wording). On top of a proof of type-safety, the paper contains a proof of operational correctness of the typed closure conversion algorithm by logical relations.

Another notable paper about closure conversion is "Typed Closure Conversion Preserves Observational Equivalence" by Ahmed and Blum [?]. The paper's title explains its main result, so we should explain the title.

(TODO bring up the Reynolds' paper) Within a language L , we have a program $P = C[A]$, where A is an implementation of an abstraction and C is the "context", or "the rest of the program". Given some other implementation A' of the abstraction, we say that A and A' are contextually equivalent when for all possible contexts C , programs $P = C[A]$ and $P' = C[A']$ behave identically.

We say that another abstraction A' is contextually equivalent to A if for all contexts C , programs $C[A]$ and $C[A']$ are equivalent. This corresponds to a programmer's intuition that A and A' behave in the same way in all possible programs.

TODO OE matters for security and safety: If an attack would be possible by exposing a certain implementation detail, then this detail is made inaccessible / private, for example by using an existential type.

Why this matters: modern software systems are made up of multiple components, of which some might not be trusted.

// To ensure reliable and secure operation, it is important to defend against faulty or malicious code. Language-based security is built upon the concept of abstraction: if access to some private implementation detail might enable an attack, then this detail is made inaccessible by hiding it behind an abstract interface, for example using an existential type. //

TODO I have quite a bit about the paper and we don't want to duplicate the paper's introduction: how do I make it shorter?

Chapter 3

Background

This chapter will introduce the relevant concepts. It will start with closure conversion, then discuss compilation phases and intermediate languages, and finally explain the Agda definitions and encodings which were borrowed from ACMM and PLFA.

3.1 Closure conversion

TODO explain and give an example

TODO explain why existential types

3.2 Compilation phases and intermediate representations

In all but the most trivial compilers, compilation proceeds in phases, or transformations. A compilation phase transforms the compilation unit to bring it one step closer from the source code to the target representation.

[diagram here]

Intermediate representations As illustrated in the figure, each compilation phase takes a source representation to a target representation [relate to diagram]. An intermediate representation can also be called an intermediate language, and abbreviated to IR or IL. For some phases, the source and target representation may be the same. Arguably, this is the case for constant expression folding.

However, other phases benefit from using different source and target representations. An example of such transformation is closure conversion, which as the reader may recall from [section], transforms abstractions with free variables to so called closures,

which take an explicit environment and can only reference values from that environment.

Typed and untyped IRs To question of whether closure conversion must necessarily use different source and target languages hinges on the distinction between typed and untyped intermediate language. Using a typed IR requires that at each point along the compilation pipeline, intermediate representations are well-typed.

Suppose that closure conversion is performed on simply typed lambda calculus (STLC). One of the two is necessary for the target language of closure conversion: either it should have first-class closures, or existential types. Neither is true of STLC, so another intermediate language is needed.

[TODO unintelligible comment about this paragraph] On the other hand, if the source and target representations are untyped, then the compiler architect might get away with using the same intermediate language as both source and target (for example Scheme, which is sometimes used as a compilation target). But even in this case, compilation process might benefit if the abstract syntax has explicit closures.

IRs in this project This project uses a dependently typed meta language (Agda) to implement compilation phases (specifically, closure conversion), so typed intermediate representations are a natural choice. Therefore, in the following sections, we will describe two intermediate representations, which are both variants of lambda calculus. The source representation will be simply typed lambda calculus, which we will refer to as λ_{st} . The target representation will be simply typed lambda calculus with closures, denoted with λ_{cl} .

The two intermediate representations are similar, and differ mainly in having either abstractions with free variables in λ_{st} or closures with environments in λ_{cl} . Unfortunately, this means that formalisations of λ_{st} and λ_{cl} share a lot of duplication. This is a common problem in formalising languages which has recently been addressed by [2]. Whether techniques from Allais et al. are applicable to this work will be discussed in [related work]. On the other hand, [section] demonstrates that while two intermediate languages can only differ in a handful of syntactic constructs and reduction steps, they can behave very differently with respect to the ubiquitous operations of renaming and substitution.

3.3 Type- and scope-safe representation of simply typed lambda calculus λ_{st}

This section will discuss the encoding of simply typed lambda calculus (abbreviated as STLC, denoted with λ_{st}), which is the source language of closure conversion. Typing and reduction rules are standard for call-by-value lambda calculus, so it is the encoding

in Agda which is of interest in this section. As similar encoding is used for the closure language λcl .

Using dependently typed Agda as the meta language allows us to encode certain invariants in the representation. Two such invariants are scope and type safety. The representation is scope-safe in the sense that all variables in a term are either bound by some binder in the term, or explicitly accounted for in the context. It is type-safe in the sense that terms are synonymous with their typing derivations, which makes ill-typed terms unrepresentable. This kind of scope and type safety is due to [4]. The rest of this section shows how this is achieved in Agda; the Agda encoding is based on the one used in [3], [2], and [9].

TODO STLC as a figure here

To start with, λst typed are defined as follows.

```
data Type : Set where
   $\alpha$       : Type
   $\_ \Rightarrow \_$  : Type  $\rightarrow$  Type  $\rightarrow$  Type
```

The context is simply a list of types.

```
Context : Set
Context = List Type
```

Variables are synonymous with proofs of context membership. Since a variable is identified by its position in the context, it is appropriate to call it a de Bruijn variable. Accordingly, the constructors of Var are named after *zero* and *successor*. Notice that the definition assumes that the leftmost type in the context corresponds to the most recently bound variable.

```
data Var : Type  $\rightarrow$  Context  $\rightarrow$  Set where
  z :  $\forall \{\sigma \Gamma\}$   $\rightarrow$  Var  $\sigma$  ( $\sigma :: \Gamma$ )
  s :  $\forall \{\sigma \tau \Gamma\}$   $\rightarrow$  Var  $\sigma$   $\Gamma$   $\rightarrow$  Var  $\sigma$  ( $\tau :: \Gamma$ )
```

We can now present the formulation of λst terms, which is synonymous with their typing derivations:

```
data Lam : Type  $\rightarrow$  Context  $\rightarrow$  Set where
  V :  $\forall \{\Gamma \sigma\}$   $\rightarrow$  Var  $\sigma$   $\Gamma$   $\rightarrow$  Lam  $\sigma$   $\Gamma$ 
  A :  $\forall \{\Gamma \sigma \tau\}$   $\rightarrow$  Lam ( $\sigma \Rightarrow \tau$ )  $\Gamma$   $\rightarrow$  Lam  $\sigma$   $\Gamma$   $\rightarrow$  Lam  $\tau$   $\Gamma$ 
  L :  $\forall \{\Gamma \sigma \tau\}$   $\rightarrow$  Lam  $\tau$  ( $\sigma :: \Gamma$ )  $\rightarrow$  Lam ( $\sigma \Rightarrow \tau$ )  $\Gamma$ 
```

The syntactic variable V constructor takes a de Bruijn variable to a term. The abstraction constructor L requires that the body is well-typed in the context Γ extended with the type σ of the variable bound by the abstraction. The application constructor A follows the typing rule for application.

3.4 Type- and scope-safe programs

Many useful traversals of the abstract syntax tree involve maintaining a mapping from free variables to appropriate values. Two such traversals are simultaneous renaming and substitution.

Simultaneous renaming takes a term N in the context Γ . It maintains a mapping ρ from variables in the original context Γ to *variables* in some other context Δ . It produces a term in Δ , which is N with variables renamed with ρ .

Similarly, simultaneous substitution takes a term N in the context Γ . It maintains a mapping σ from variables in the original context Γ to *terms* in some other context Δ . It produces a term in Δ , which is N with variables substitution for with σ .

Before we can demonstrate an implementation of renaming and substitution, we need to formalise the notion of a mapping from free variables to appropriate values, which we call the *environment*.

```
record _-Env (Γ : Context) (V : Type → Context → Set) (Δ : Context) : Set where
  constructor pack
  field lookup : ∀ → Var σ Γ → V σ Δ
```

An environment $(\Gamma \text{--Env}) \mathcal{V} \Delta$ encapsulates a mapping from variables in Γ to values \mathcal{V} (variables for renaming, terms for substitution) which are well-typed and -scoped in Δ .

An environment which maps variables to variables is important enough to deserve its own name.

```
Thinning : Context → Context → Set
Thinning Γ Δ = (Γ -Env) Var Δ
```

There is a notion of an empty environment ε , of extending an environment ρ with a value v : $\rho \bullet v$, and of mapping a function f over an environment ρ : $f \text{<}\$ \text{>} \rho$, corresponding to the analogous operations on contexts (which are just lists). Finally, `select ren ρ` renames a variable with `ren` before looking it up in ρ .

```
ε : ∀ {V Δ} → ([] -Env) V Δ
lookup ε ()

_•_ : ∀ {Γ Δ σ V} → (Γ -Env) V Δ → V σ Δ → (σ :: Γ -Env) V Δ
lookup (ρ • v) Z = v
lookup (ρ • v) (S x) = lookup ρ x

_<$>_ : ∀ {Γ Δ Θ V1 V2}
        → (V → V1 σ Δ → V2 σ Θ) → (Γ -Env) V1 Δ → (Γ -Env) V2 Θ
lookup (f <$> ρ) x = f (lookup ρ x)

select : ∀ {Γ Δ Θ V} → Thinning Γ Δ → (Δ -Env) V Θ → (Γ -Env) V Θ
lookup (select ren ρ) k = lookup ρ (lookup ren k)
```

Notice that those four operations on environments are defined using copatterns [1] by ‘observing’ the behaviour of lookup.

Equipped with the notion of environments, we can give an implementation of renaming and substitution:

```

ext : ∀ {Γ Δ} {σ : Type} → Thinning Γ Δ → Thinning (σ :: Γ) (σ :: Δ)
ext ρ = s <$> ρ • z

rename : ∀ {Γ Δ σ} → Thinning Γ Δ → Lam σ Γ → Lam σ Δ
rename ρ (V x)      = V (lookup ρ x)
rename ρ (L N)      = L (rename (ext ρ) N)
rename ρ (A M N)    = A (rename ρ M) (rename ρ N)

exts : ∀ {Γ Δ} {τ : Type} → (Γ -Env) Lam Δ → (τ :: Γ -Env) Lam (τ :: Δ)
exts σ = rename (pack s) <$> σ • V z

subst : ∀ {Γ Δ σ} → (Γ -Env) Lam Δ → Lam σ Γ → Lam σ Δ
subst σ (V x)      = lookup σ x
subst σ (L N)      = L (subst (exts σ) N)
subst σ (A M N)    = A (subst σ M) (subst σ N)

```

Notice that those two traversals are identical except (1) *renaming* wraps the result of `lookup ρ x` in `V`, and *renaming* and *substitution* extend the environment in a different way: `s <$> ρ • z` vs `rename (pack s) <$> σ • V z`. The observation that renaming and substitution for STLC share a common structure was a basis was the unpublished manuscript by McBride [5], and subsequently motivated the ACMM paper [3]. In [section], we will show how ACMM abstracts this common structure of renaming and substitution into a notion of a semantics.

Also notice how the functions `ext` and `exts` extend the environment when the traversal goes under a binder.

An example instantiation of simultaneous substitution is single substitution. Single substitution replaces occurrences of the last-bound variable in the context, and it is useful for defining the beta reduction for abstractions. Single substitution environment is an identity substitution environment extended with a single value:

```

id-subst : ∀ → Subst Γ Γ
lookup id-subst x = V x

_/_ : ∀ {Γ σ τ} → Lam τ (σ :: Γ) → Lam σ Γ → Lam τ Γ
_/_ {} N M = subst (id-subst • M) N

```

3.5 ACMM’s notion of a semantics

TODO ACMM, synch, fusions

3.6 Small-step operational semantics

The formalisation of small step semantics for a call-by-value lambda calculus is adapted from [9].

Values are terms which do not reduce further. In this most basic version of lambda calculus language, the only values are abstractions:

data Value : $\forall \{\Gamma \sigma\} \rightarrow \text{Lam } \sigma \Gamma \rightarrow \text{Set where}$

V-L : $\forall \{\Gamma \sigma \tau\} \{N : \text{Lam } \tau (\sigma :: \Gamma)\}$

$\rightarrow \text{Value } (\text{L } N)$

Our operational semantics include two kinds of reduction rules. Compatibility rules, whose names start with ξ , reduce parts of the term (specifically, the LHS and RHS of application). Beta reduction β -L, on the other hand, describes what an abstraction applied to a value reduces to.

data $_ \longrightarrow _$: $\forall \{\Gamma \sigma\} \rightarrow (\text{Lam } \sigma \Gamma) \rightarrow (\text{Lam } \sigma \Gamma) \rightarrow \text{Set where}$

ξ -A₁ : $\forall \{\Gamma \sigma \tau\} \{M M' : \text{Lam } (\sigma \Rightarrow \tau) \Gamma\} \{N : \text{Lam } \sigma \Gamma\}$

$\rightarrow M \longrightarrow M'$

$\rightarrow \text{A } M N \longrightarrow \text{A } M' N$

ξ -A₂ : $\forall \{\Gamma \sigma \tau\} \{V : \text{Lam } (\sigma \Rightarrow \tau) \Gamma\} \{N N' : \text{Lam } \sigma \Gamma\}$

$\rightarrow \text{Value } V$

$\rightarrow N \longrightarrow N'$

$\rightarrow \text{A } V N \longrightarrow \text{A } V N'$

β -L : $\forall \{\Gamma \sigma \tau\} \{N : \text{Lam } \tau (\sigma :: \Gamma)\} \{V : \text{Lam } \sigma \Gamma\}$

$\rightarrow \text{Value } V$

$\rightarrow \text{A } (\text{L } N) V \longrightarrow N / V$

A term which can take a reduction step is called a reducible expression, or a redex. A property of a language that every well-typed term is either a value or a redex is called type-safety. This property is captured by a slogan ‘well-typed terms don’t get stuck’ and can be proved by techniques like ‘progress and preservation’ or logical relations. Simply typed lambda calculus is type-safe, and so is this formalisation. For a proof of type safety for a similar formalisation of STLC, cf. [9].

Operational semantics are needed for the treatment of bisimulation.

Chapter 4

Formalising closure conversion

This chapter presents this project’s formalisation of closure conversion. It starts by discussing the closure language λcl , an intermediate language which is like STLC but with abstractions replaced by closures. Then it demonstrates a type-preserving conversion for λst to λcl which has the property that the obtained closure environments are ‘minimal’. Finally, several properties about interactions between renaming and substitution in λcl are formally established — they are needed in proofs of correctness in subsequent chapters.

4.1 Closure language λcl

As discussed in the Background [or maybe Intro?] chapter, some compilation phases must use different source and target intermediate representations. This is the case with closure conversion, and this section presents a formalisation of an intermediate language with closures. The language is very similar the formalised simply typed lambda calculus, except that abstraction with free variables are replaced by closures with environments. What might seem like a simple change has interesting implications for traversals like renaming and substitution.

The closure language λcl shares types, contexts, and de-Brujin-variables-as-proofs-of-context-membership, and their respective Agda formalisations, with the source representation. In general, two different intermediate representations do not need to share the same type system, but if they do, this simplifies formalisation. The descriptions of those formalisations can be found in Section [TODO].

4.1.1 Terms

The definition of terms of λcl differs from terms of λst in the L constructor, which, in λcl , holds the closure body and the closure environment.

```
data Lam : Type → Context → Set where
  V  : ∀ {Γ σ}      → Var σ Γ      → Lam σ Γ
```

$$\begin{array}{lcl}
\text{A} & : \forall \{\Gamma \sigma \tau\} & \rightarrow \text{Lam } (\sigma \Rightarrow \tau) \Gamma \rightarrow \text{Lam } \sigma \Gamma \rightarrow \text{Lam } \tau \Gamma \\
\text{L} & : \forall \{\Gamma \Delta \sigma \tau\} & \rightarrow \text{Lam } \tau (\sigma :: \Delta) \rightarrow (\Delta \text{ --Env}) \text{Lam } \Gamma \rightarrow \text{Lam } (\sigma \Rightarrow \tau) \Gamma
\end{array}$$

Notice that the typing rule for the closure constructor L mentions two contexts, Γ and Δ . We call Γ the *outer context* and Δ the *inner context* of a closure.

$$\frac{\Gamma, x : \sigma \vdash e : \tau}{\Gamma \vdash \lambda x : \sigma. e : \sigma \rightarrow \tau} \text{T-abs} \qquad \frac{e_{ev} = \text{subst}(\Delta \subseteq \Gamma) \quad \Delta, x : \sigma \vdash e : \tau}{\Gamma \vdash \langle\langle \lambda x : \sigma. e, e_{ev} \rangle\rangle : \sigma \rightarrow \tau} \text{T-clos}$$

The closure as a whole is typed in Γ , but the closure body (also called the *closure code*) is typed in $\sigma :: \Delta$. The relationship between Γ and Δ is given by the closure environment.

A closure environment is traditionally implemented as a record, and variables in the closure code reference fields of that record. In this development, on the other hand, the environment is represented as a substitution environment, that is, a mapping from variables in Δ to terms in Γ . This representation is isomorphic to the one using a record, and it has several benefits, especially eliminating the need for products in the language, and overall simplification of the formalisation.

Finally, recall from [section] that in order for a closure-converted program to be well-typed, a closure environment should have an existential type. It is important to note that in this formalisation, existential typing is achieved in the meta language Agda, not in the object language λcl , which does not have existential types. Indeed, existential quantification (including over types) can be achieved in Agda through dependent products, a datatype constructor is a dependent product, and the environment is a parameter to the L constructor.

4.1.2 Renaming and substitution

Consider the case for the constructor L of renaming and substitution in λcl and how it is different from the corresponding definition in λst .

$$\begin{array}{lcl}
\text{rename} & : \forall \{\Gamma \Delta \sigma\} & \rightarrow \text{Thinning } \Gamma \Delta \rightarrow \text{Lam } \sigma \Gamma \rightarrow \text{Lam } \sigma \Delta \\
\text{rename } \rho & (\text{V } x) & = \text{V } (\text{lookup } \rho x) \\
\text{rename } \rho & (\text{A } M N) & = \text{A } (\text{rename } \rho M) (\text{rename } \rho N) \\
\text{rename } \rho & (\text{L } N E) & = \text{L } N (\text{rename } \rho \text{ <\$> } E)
\end{array}$$

$$\begin{array}{lcl}
\text{subst} & : \forall \{\Gamma \Delta \sigma\} & \rightarrow \text{Subst } \Gamma \Delta \rightarrow \text{Lam } \sigma \Gamma \rightarrow \text{Lam } \sigma \Delta \\
\text{subst } \rho & (\text{V } x) & = \text{lookup } \rho x \\
\text{subst } \rho & (\text{A } M N) & = \text{A } (\text{subst } \rho M) (\text{subst } \rho N) \\
\text{subst } \rho & (\text{L } N E) & = \text{L } N (\text{subst } \rho \text{ <\$> } E)
\end{array}$$

Unlike in λst , renaming and substitution in λcl *do not go under binders* (do not change the closure body). This is because renaming and substitution take a term in a context Γ to a term in a context Γ' . But the code (body) of a closure is typed in a different

context Δ . So upon recursing on a closure, renaming and substitution adjust the closure environment and leave the closure body unchanged. The adjustment to the environment is $\text{rename } \rho \langle \$ \rangle E$ in the case of renaming and $\text{subst } \rho \langle \$ \rangle E$ in the case of substitution. In either case, the adjustment consists of mapping the renaming/substitution over the values in the environment.

The fact that in λcl , renaming and substitution do not go under binders will allow us to prove ‘fusion lemmas’ in [section] without using the machinery of ACMM, which will significantly simplify the proofs.

Just like in λst , we also define functions ext and exts which extend the environment when renaming or substitution goes under a binder:

$$\begin{aligned} \text{ext} &: \forall \{\Gamma \Delta\} \{\sigma : \text{Type}\} \rightarrow \text{Thinning } \Gamma \Delta \rightarrow \text{Thinning } (\sigma :: \Gamma) (\sigma :: \Delta) \\ \text{ext } \rho &= \text{pack } s \langle \$ \rangle \rho \bullet z \end{aligned}$$

$$\begin{aligned} \text{exts} &: \forall \{\Gamma \Delta \sigma\} \rightarrow \text{Subst } \Gamma \Delta \rightarrow \text{Subst } (\sigma :: \Gamma) (\sigma :: \Delta) \\ \text{exts } \rho &= \text{rename } (\text{pack } s) \langle \$ \rangle \rho \bullet V z \end{aligned}$$

4.1.3 Operational semantics

Operational semantics are similar to the semantics for λst , except for adjustments for closures. Values in λcl are closures, and the rule for beta reduction is different:

$$\begin{aligned} &\text{infix 2 } _ \longrightarrow _ \\ &\text{data } _ \longrightarrow _ : \forall \{\Gamma \sigma\} \rightarrow (\text{Lam } \sigma \Gamma) \rightarrow (\text{Lam } \sigma \Gamma) \rightarrow \text{Set where} \\ &\quad \beta\text{-L} : \forall \{\Gamma \Delta \sigma \tau\} \{N : \text{Lam } \tau (\sigma :: \Delta)\} \{E : \text{Subst } \Delta \Gamma\} \{V : \text{Lam } \sigma \Gamma\} \\ &\quad \rightarrow \text{Value } V \\ &\quad \rightarrow \text{A } (\text{L } N E) V \longrightarrow \text{subst } (E \bullet V) N \end{aligned}$$

Recall that a closure is a function without free variables, partially applied to an environment. When the closure argument reduces to a value, the argument and the values in the environment get simultaneously substituted into the closure body. The simplicity of this reduction rule is another benefit of representing environments as substitution environments.

4.1.4 Conversion from λst to λcl

This project’s approach to typed, or type-preserving, closure conversion follows [7]. An important point here is that the specification of typed closure conversion allows for different implementations which might differ in their treatment of environments. The only requirement in the specification is that

1. If the source term is an abstraction typed in the context Γ ;
2. if the body of the source abstraction can be typed in a smaller context Δ , such that $\Delta \subseteq \Gamma$;
3. then the target terms is a closure whose environment is a substitution from Δ to Γ .

This is given by the following conversion rule:

$$\frac{e_{ev} = subst(\Delta \subseteq \Gamma) \quad \Delta, x : \sigma \vdash e \rightsquigarrow e' : \tau}{\Gamma \vdash \lambda x : \sigma. e \rightsquigarrow \langle \langle \lambda x : \sigma. e', e_{ev} \rangle \rangle : \sigma \rightarrow \tau}$$

It is up to the implementation of closure conversion to decide how big to make Δ , on the spectrum between (1) Δ being equal to Γ , and (2) Δ being ‘minimal’, i.e. only containing the parts of Γ which are necessary to type the term. We present two Agda implementation of closure conversion, corresponding to the two ends of the spectrum.

Closure conversion where Δ is the same as Γ is a simple transformation:

```
simple-cc : ∀ {Γ σ} → S.Lam σ Γ → T.Lam σ Γ
simple-cc (S.V x) = T.V x
simple-cc (S.A M N) = T.A (simple-cc M) (simple-cc N)
simple-cc (S.L N) = T.L (simple-cc N) T.id-subst
```

where `T.id-subst` is the identity substitution which maps a term in Γ to itself, defined as:

```
id-subst : ∀ → Subst Γ Γ
lookup id-subst x = V x
```

We call the other end of the spectrum *minimising closure conversion*. Its implementation in Agda is rather more involved and is described in the next section.

4.1.5 Minimising closure conversion

Minimising closure conversion is given by the following deduction rules, where a statement $\Gamma \vdash e : \sigma \rightsquigarrow \Delta \vdash e' : \sigma$ should be read as: ‘the term e of type σ in the context Γ can be closure converted to the term e' in Δ ’:

$$\frac{}{\Gamma \vdash x : \sigma \rightsquigarrow \emptyset, x : \sigma \vdash x : \sigma} \text{ (min-V)} \quad \frac{\Gamma \vdash e_1 : \sigma \rightarrow \tau \rightsquigarrow \Delta_1 \vdash e'_1 : \sigma \rightarrow \tau \quad \Gamma \vdash e_2 : \sigma \rightsquigarrow \Delta_2 \vdash e'_2 : \sigma \quad \Delta = merge \Delta_1 \Delta_2}{\Gamma \vdash e_1 e_2 : \tau \rightsquigarrow \Delta \vdash e'_1 e'_2 : \tau} \text{ (min-A)}$$

$$\frac{\Gamma, x : \sigma \vdash e : \tau \rightsquigarrow \Delta, x : \tau \vdash e : \tau \quad e_{id} = subst(\Delta \subseteq \Gamma)}{\Gamma \vdash \lambda x : \sigma. e : \sigma \rightarrow \tau \rightsquigarrow \Delta \vdash \langle \langle \lambda x : \sigma. e, e_{id} \rangle \rangle : \sigma \rightarrow \tau} \text{ (min-L)}$$

min-V: Any variables can be typed in a singleton context containing just the type of the variable.

min-A: If the conversion e_1' of e_1 can be typed in Δ_1 , and the conversion e_2' of e_2 can be typed in Δ_2 , then the application $e_1' e_2'$ can be typed in Δ , where Δ is the result of merging Δ_1 and Δ_2 .

min-L: If the conversion e' of the abstraction body e can be typed in context $\sigma :: \Delta$ (or $\Delta, x : \sigma$, using the notation with names), then the closure resulting from the conversion of the abstraction can be typed in Δ , and it has the identity environment $\Delta \subseteq \Delta$.

To formalise this conversion in Agda, we need several helper definitions.

4.1.5.1 Merging subcontexts

The deduction rules for minimising closure conversion contained statements of the form $\Delta \subseteq \Gamma$, which reads: ‘ Δ is a subcontext of Γ ’. Since in this development, a context is just a list of types, the notion of subcontexts can be captured with the `_⊆_` (sublist) relation from Agda’s standard library. The inductive definition of the relation is:

```
data _⊆_ : List A → List A → Set where
  base   : [] ⊆ []
  skip   : ∀ {xs y ys} → xs ⊆ ys → xs ⊆ (y :: ys)
  keep   : ∀ {x xs ys} → xs ⊆ ys → (x :: xs) ⊆ (x :: ys)
```

This project’s contribution is to define the operation of merging two subcontexts. Given contexts Γ , Δ , and Δ_1 such that $\Delta \subseteq \Gamma$ and $\Delta_1 \subseteq \Gamma$, the result of merging the subcontexts Δ and Δ_1 is a context Γ_1 which satisfies the following conditions:

1. It is contained in the big context: $\Gamma_1 \subseteq \Gamma$.
2. It contains the small contexts: $\Delta \subseteq \Gamma_1$ and $\Delta_1 \subseteq \Gamma_1$.
3. The proof that $\Delta \subseteq \Gamma$ obtained by transitivity from $\Delta \subseteq \Gamma_1$ and $\Gamma_1 \subseteq \Gamma$ is the same as the input proof that $\Delta \subseteq \Gamma$; similarly for $\Delta_1 \subseteq \Gamma$.

All those requirements are captured by the following dependent record in Agda:

```
record SubListSum {Γ Δ Δ₁ : List A} (Δ⊆Γ : Δ ⊆ Γ) (Δ₁⊆Γ : Δ₁ ⊆ Γ) : Set where
  constructor subListSum
  field
    Γ₁      : List A
    Γ₁⊆Γ    : Γ₁ ⊆ Γ
    Δ⊆Γ₁    : Δ ⊆ Γ₁
    Δ₁⊆Γ₁   : Δ₁ ⊆ Γ₁
    well    : ⊆-trans Δ⊆Γ₁ Γ₁⊆Γ ≡ Δ⊆Γ
    well₁   : ⊆-trans Δ₁⊆Γ₁ Γ₁⊆Γ ≡ Δ₁⊆Γ
```

The type of the function which merges two subcontexts can be stated as:

$$\text{merge} : \forall \{ \Gamma \Delta \Delta_1 \} \rightarrow (\Delta \subseteq \Gamma : \Delta \subseteq \Gamma) \rightarrow (\Delta_1 \subseteq \Gamma : \Delta_1 \subseteq \Gamma) \rightarrow \text{SubListSum } \Delta \subseteq \Gamma \Delta_1 \subseteq \Gamma$$

We argue that the type of the function completely captures its behaviour (TODO how would we prove this?). The fact that a type can completely capture the behaviour of a function is a remarkable feature of programming with dependent types. Even more remarkable is the fact that the logical properties of Γ_1 are useful computationally. E.g the proof that $\Delta \subseteq \Gamma_1$ determines a renaming from Δ to Γ_1 , which is used in the minimising closure conversion algorithm. A further example: the fact that \subseteq -trans $\Delta \subseteq \Gamma_1 \Gamma_1 \subseteq \Gamma \equiv \Delta \subseteq \Gamma$ is used in proofs of certain equivalences involving subcontexts and renaming.

4.1.6 Agda implementation of minimising closure conversion

Recall that terms of our intermediate languages are explicitly typed in a given context. For that reason, the result type of minimising closure conversion must be existentially quantified over a context. In fact, the context should be a subcontext of the input context Γ . This is captured with the dependent record $_ \vdash _$:

```
record _⊢_ (Γ : Context) (A : Type) : Set where
  constructor ∃[ ] ^_
  field
    Δ : Context
    Δ⊆Γ : Δ ⊆ Γ
    N : T.Lam A Δ
```

For example, a term N in a context Δ which is a subcontext of Γ by $\Delta \subseteq \Gamma$, would be constructed as $\exists[\Delta] \Delta \subseteq \Gamma \wedge N$.

With this data type, the type of the minimising closure conversion function is:

$$\text{cc} : \forall \{ \Gamma A \} \rightarrow \text{S.Lam } A \Gamma \rightarrow \Gamma \vdash A$$

The function definition is by cases:

Variable case

$$\text{cc } \{ A = A \} (\text{S.V } x) = \exists[A :: []] \text{Var} \rightarrow \subseteq x \wedge \text{T.V } z$$

Following *min-V*, a variable is typed in a singleton context. The proof of the subcontext relation is computed from the proof of the context membership by a function $\text{Var} \rightarrow \subseteq$.

Application case

$$\begin{aligned} & \text{cc } (\text{S.A } M N) \text{ with } \text{cc } M \mid \text{cc } N \\ & \text{cc } (\text{S.A } M N) \mid \exists[\Delta] \Delta \subseteq \Gamma \wedge M^\dagger \mid \exists[\Delta_1] \Delta_1 \subseteq \Gamma \wedge N^\dagger \text{ with } \text{merge } \Delta \subseteq \Gamma \Delta_1 \subseteq \Gamma \\ & \text{cc } (\text{S.A } M N) \mid \exists[\Delta] \Delta \subseteq \Gamma \wedge M^\dagger \mid \exists[\Delta_1] \Delta_1 \subseteq \Gamma \wedge N^\dagger \mid \text{subListSum } \Gamma_1 \Gamma_1 \subseteq \Gamma \Delta \subseteq \Gamma_1 \Delta_1 \subseteq \Gamma_1 \dots \\ & = \exists[\Gamma_1] \Gamma_1 \subseteq \Gamma \wedge (\text{T.A } (\text{T.rename } (\subseteq \rightarrow \rho \Delta \subseteq \Gamma_1) M^\dagger) (\text{T.rename } (\subseteq \rightarrow \rho \Delta_1 \subseteq \Gamma_1) N^\dagger)) \end{aligned}$$

Given an application $e_1 e_2$, e_1 and e_2 are closure converted recursively, resulting in terms e_1' and e_2' , which are typed in Δ_1 and Δ_2 , respectively. Following *app-V*, the result of closure-converting the application is typed in the context Δ , which is the result of merging Δ_1 and Δ_2 . As terms are explicitly typed in a context, e_1' and e_2' have to be renamed from Δ_1 to Δ , and from Δ_2 to Δ , respectively. A renaming environment is computed from a subcontext relation proof by the function $\subseteq \rightarrow \rho$ which is given by:

```

 $\subseteq \rightarrow \rho : \{\Gamma \Delta : \text{Context}\} \rightarrow \Gamma \subseteq \Delta \rightarrow \text{Thinning } \Gamma \Delta$ 
lookup ( $\subseteq \rightarrow \rho$  base) ()
lookup ( $\subseteq \rightarrow \rho$  (skip  $\Gamma \subseteq \Delta$ ))  $x = s$  (lookup ( $\subseteq \rightarrow \rho$   $\Gamma \subseteq \Delta$ )  $x$ )
lookup ( $\subseteq \rightarrow \rho$  (keep  $\Gamma \subseteq \Delta$ ))  $z = z$ 
lookup ( $\subseteq \rightarrow \rho$  (keep  $\Gamma \subseteq \Delta$ )) ( $s x$ ) =  $s$  (lookup ( $\subseteq \rightarrow \rho$   $\Gamma \subseteq \Delta$ )  $x$ )

```

Abstraction case

```

cc (S.L N) with cc N
cc (S.L N) |  $\exists [\Delta] \Delta \subseteq \Gamma \wedge N^\dagger$  with adjust-context  $\Delta \subseteq \Gamma$ 
cc (S.L N) |  $\exists [\Delta] \Delta \subseteq \Gamma \wedge N^\dagger$  | adjust  $\Delta_1 \Delta_1 \subseteq \Gamma \Delta \subseteq A \Delta_1$  _
=  $\exists [\Delta_1] \Delta_1 \subseteq \Gamma \wedge (\text{T.L } (\text{T.rename } (\subseteq \rightarrow \rho \Delta \subseteq A \Delta_1) N^\dagger) \text{T.id-subst})$ 

```

Following *min-A*, the result of closure-converting an abstraction depends on the result N^\dagger of closure-clonverting its body. A recursive call on the body of the abstraction yields a term typed in some context Δ . But looking at the typing rule for closures (*T-clos*), the closure body is typed in a context $\sigma :: \Delta_1$ (or $\Delta_1, x : \sigma$ using named variables), where σ is the type of the last bound variable and Δ_1 is the context corresponding to the closure environment. Thus, we need a way of decomposing Δ into σ and Δ_1 , together with an appropriate proof of membership in the input context Γ .

This task is achieved by the function *adjust-context*:

```

adjust-context :  $\forall \{\Gamma \Delta A\} \rightarrow (\Delta \subseteq A :: \Gamma : \Delta \subseteq A :: \Gamma) \rightarrow \text{AdjustContext } \Delta \subseteq A :: \Gamma$ 

```

whose specification is captured by its return type which uses the dependent record *AdjustContext*:

```

record AdjustContext {A  $\Gamma \Delta$ } ( $\Delta \subseteq A :: \Gamma : \Delta \subseteq A :: \Gamma$ ) : Set where
  constructor adjust
  field
     $\Delta_1$       : Context
     $\Delta_1 \subseteq \Gamma$  :  $\Delta_1 \subseteq \Gamma$ 
     $\Delta \subseteq A \Delta_1$  :  $\Delta \subseteq A :: \Delta_1$ 
    well      :  $\Delta \subseteq A :: \Gamma \equiv \subseteq\text{-trans } \Delta \subseteq A \Delta_1$  (keep  $\Delta_1 \subseteq \Gamma$ )

```

The specification is: given $\Delta \subseteq A :: \Gamma$, there exists a context Δ_1 such that $\Delta_1 \subseteq \Gamma$ and $\Delta \subseteq A :: \Delta_1$, such that the proof $\Delta \subseteq A :: \Gamma$ obtained by transitivity is the same as the input proof.

The evidence that $\Delta \subseteq A :: \Delta_1$ is used to rename N^\dagger so that the final inherently-typed term is well-typed.

We also provide a wrapper function $_^\dagger$:

$$\begin{aligned} _^\dagger &: \forall \{ \Gamma A \} \rightarrow \text{S.Lam } A \Gamma \rightarrow \text{T.Lam } A \Gamma \\ _^\dagger M &\text{ with cc } M \\ _^\dagger M \mid \exists! [\Delta] \Delta \subseteq \Gamma \wedge N &= \text{T.rename } (\subseteq \rightarrow \rho \Delta \subseteq \Gamma) N \end{aligned}$$

This function is a wrapper over the min-cc function which undoes the minimisation on the outer level. In other words, all closures in the term are still minimised, but the outer term is typed in the same context as the input source term. This is useful when we need to compare the input and output of closure conversion, and need to ensure that they are typed in the same context.

4.1.7 Fusion lemmas for the closure language λcl

When studying the meta-theory of a calculus, one systematically needs to prove fusion lemmas for various traversals. A fusion lemma relates three traversals: the pair we sequence and their sequential composition. The two traversals which have to be fused in later proofs are renaming and substitution. There are four ways we can sequence renaming and substitution, and each of those four sequencing can be expressed as a single renaming or substitution:

1. A renaming followed by a renaming,
2. A renaming followed by a substitution,
3. A substitution followed by a renaming,
4. A substitution followed by a substitution.

We state the results as signatures of Agda functions, using the environment combinators $_<\$>_$ and select which are described in Section 3.4.

$$\begin{aligned} \text{rename} \circ \text{rename} &: \forall \{ \Gamma \Delta \Theta \tau \} (\rho_1 : \text{Thinning } \Gamma \Delta) (\rho_2 : \text{Thinning } \Delta \Theta) (N : \text{Lam } \tau \Gamma) \\ &\rightarrow \text{rename } \rho_2 (\text{rename } \rho_1 N) \equiv \text{rename } (\text{select } \rho_1 \rho_2) N \end{aligned}$$

$$\begin{aligned} \text{subst} \circ \text{rename} &: \forall \{ \Gamma \Delta \Theta \tau \} (\rho\sigma : \text{Subst } \Gamma \Theta) (\rho\rho : \text{Thinning } \Delta \Gamma) (N : \text{Lam } \tau \Delta) \\ &\rightarrow \text{subst } \rho\sigma (\text{rename } \rho\rho N) \equiv \text{subst } (\text{select } \rho\rho \rho\sigma) N \end{aligned}$$

$$\begin{aligned} \text{rename} \circ \text{subst} &: \forall \{ \Gamma \Delta \Theta \tau \} (\rho\rho : \text{Thinning } \Gamma \Theta) (\rho\sigma : \text{Subst } \Delta \Gamma) (N : \text{Lam } \tau \Delta) \\ &\rightarrow \text{rename } \rho\rho (\text{subst } \rho\sigma N) \equiv \text{subst } (\text{rename } \rho\rho _<\$> \rho\sigma) N \end{aligned}$$

$$\begin{aligned} \text{subst} \circ \text{subst} &: \forall \{ \Gamma \Delta \Theta \tau \} (\rho_1 : \text{Subst } \Gamma \Theta) (\rho_2 : \text{Subst } \Delta \Gamma) (N : \text{Lam } \tau \Delta) \\ &\rightarrow \text{subst } \rho_1 (\text{subst } \rho_2 N) \equiv \text{subst } (\text{subst } \rho_1 _<\$> \rho_2) N \end{aligned}$$

Rather than include Agda proofs of all four lemmas, here we outline the proof structure, analyse just one of the four proofs, and compare fusion lemmas for λ_{cl} with the corresponding lemmas for λ_{st} .

A generic technique to prove fusion lemmas for STLC, including the ones about renaming and substitution, is one of the main contributions of ACMM [3]. Their proof uses Kripke logical relations and it relies on the invariant that corresponding environment values are in appropriate relations, including when environments are extended when going under a binder. Maintaining this invariant is possible thanks to the generic framework for writing traversals introduced by ACMM.

As it turns out, fusion lemmas for the closure language are simpler, as they do not require the logical relation machinery of ACMM. This is because renaming and substitution in λ_{cl} ‘do not go under binders’, as can be seen from their definitions in Section 4.1.2. For both renaming and substitution, in the closure case (L), the closure body is left untouched; only the closure environment is modified.

We are now ready to take a closer look at the proof of the fusion lemma stating that a renaming followed by a substitution is a substitution:

$$\begin{aligned}
& \text{subst}\circ\text{rename} : \forall \{\Gamma \Delta \Theta \tau\} (\rho\sigma : \text{Subst } \Gamma \Theta) (\rho\rho : \text{Thinning } \Delta \Gamma) (N : \text{Lam } \tau \Delta) \\
& \quad \rightarrow \text{subst } \rho\sigma (\text{rename } \rho\rho N) \equiv \text{subst } (\text{select } \rho\rho \rho\sigma) N \\
\\
& \text{subst}\circ\text{rename } \rho\sigma \rho\rho (\mathbf{V} \ x) = \text{refl} \\
& \text{subst}\circ\text{rename } \rho\sigma \rho\rho (\mathbf{A} \ M \ N) = \text{cong}_2 \ \mathbf{A} \ (\text{subst}\circ\text{rename } \rho\sigma \rho\rho M) \\
& \quad (\text{subst}\circ\text{rename } \rho\sigma \rho\rho N) \\
& \text{subst}\circ\text{rename } \rho\sigma \rho\rho (\mathbf{L} \ N \ E) = \text{cong}_2 \ \mathbf{L} \ \text{refl} \ (\text{env-extensionality } h) \\
& \text{where } h : (_<\$>_ \{ \mathcal{W} = \text{Lam} \} (\text{subst } \rho\sigma) (_<\$>_ \{ \mathcal{W} = \text{Lam} \} (\text{rename } \rho\rho) E)) \\
& \quad \equiv^E (\text{subst } (\text{select } \rho\rho \rho\sigma) <\$> E) \\
& \quad h = \text{begin}^E \\
& \quad \quad _<\$>_ \{ \mathcal{W} = \text{Lam} \} (\text{subst } \rho\sigma) (_<\$>_ \{ \mathcal{W} = \text{Lam} \} (\text{rename } \rho\rho) E) \\
& \quad \equiv^E \langle _<\$>\text{-distr } \{ \mathcal{W} = \text{Lam} \} (\text{rename } \rho\rho) (\text{subst } \rho\sigma) E \rangle \\
& \quad \equiv^E \langle _<\$>_ \{ \mathcal{W} = \text{Lam} \} (\text{subst } \rho\sigma \circ \text{rename } \rho\rho) E \rangle \\
& \quad \equiv^E \langle _<\$>\text{-fun } \{ \mathcal{W} = \text{Lam} \} (\lambda e \rightarrow \text{subst}\circ\text{rename } \rho\sigma \rho\rho e) E \rangle \\
& \quad \quad \text{subst } (\text{select } \rho\rho \rho\sigma) <\$> E \\
& \quad \blacksquare^E
\end{aligned}$$

The proof is by induction on the typing derivation of the term:

- In the variable case, the LHS and the RHS normalise to the same term, so `refl` suffices.
- In the application case, the proof is by induction and congruence.
- In the closure case, the proof is also by congruence, but an equational proof is required to show that the LHS and RHS act in the same way on the environment E .

The equational proof proceeds as follows:

1. It uses the fact that function composition $_ \circ _$ distributes through mapping over environments $_ \langle \$ \rangle _$: we have $f \langle \$ \rangle g \langle \$ \rangle E \equiv f \circ g \langle \$ \rangle E$ which is captured by the lemma $\langle \$ \rangle$ -distr,
2. It uses the fact that when f and g are extensionally equal ($\forall \{x\} \rightarrow f\ x \equiv g\ x$), then $f \langle \$ \rangle E \equiv g \langle \$ \rangle E$ which is captured by the lemma $\langle \$ \rangle$ -fun,
3. $\langle \$ \rangle$ -fun is instantiated with the inductive hypothesis.

Unfortunately, Agda does not recognise this project's fusion lemmas as terminating, and we were unable to provide a termination proof. Still, we believe that the function does in fact terminate.

Chapter 5

Proving correctness of closure conversion with bisimulation

Preceding sections defined the source and target languages of closure conversion, λ_{st} and λ_{cl} , together with reduction rules for each, and a closure conversion function min-cc from λ_{st} to λ_{cl} .

The min-cc closure conversion is type- and scope-preserving by construction. The property of type preservation provides confidence in the compilation process, but in this theoretical development which deals with a small, toy language, it is within the reach of this project to prove properties about operational correctness.

One such operational correctness property of a pair of languages is **bisimulation**. Intuition about bisimulation is captured by a slogan: similar terms reduce to similar terms.

This chapter starts by defining a relation between terms of λ_{st} and terms of λ_{cl} , which we call a *compatibility relation*. The compatibility relation is syntactic: in general, two terms are compatible when their subterms are compatible.

Then, we define what it means for a relation to be a bisimulation. A bisimulation is a relation which has a semantic property which relates reduction steps of source and target terms. Next, we will show that the compatibility relation is a bisimulation.

Finally, we will link the compatibility relation to closure conversion: we will argue that the graph relation of every sensible closure conversion function is contained in the compatibility relation. In particular, we will prove that this is the case for min-cc .

Overall, correctness of the minimising closure conversion is established: first, by showing that the input and output of closure conversion are related by a syntactic relation, and second, by showing that this syntactic relation is also a semantic relation. Thus, soundness of our closure conversion is established.

The part which shows that the compatibility relation is a bisimulation is inspired by the ‘Bisimulation’ chapter from [9].

5.0.1 Compatibility relation

The compatibility relation is defined as follows:

Definition. Given a term M in λ_{st} and a term M^\dagger in λ_{cl} , the compatibility relation $M \sim M^\dagger$ is defined inductively as follows:

- (*Variable*) For any given variable (proof of context membership) x , we have $S.\text{' } x \sim T.\text{' } x$.
- (*Application*) If $M \sim M^\dagger$ and $N \sim N^\dagger$, then $M \cdot N \sim M^\dagger \cdot N^\dagger$.
- (*Abstraction*) If $N \sim T.\text{subst } (T.\text{exts } E) N^\dagger$, then $S.L N \sim T.L N^\dagger E$.

Recall that λ_{st} and λ_{cl} share types, contexts, and variables (proofs of context membership). In fact, compatibility is only defined for source and target terms of the same type in the same context (this is explicit in the Agda definition).

While the variable and application cases are straightforward, the abstraction / closure case needs some explanation. Since the body N of the abstraction is defined in $\sigma :: \Gamma$, and the body of the closure N^\dagger is defined in $\sigma :: \Delta$, they cannot be compatible. However, N can be compatible with the result of substituting the environment E in N^\dagger (the environment is extended with a variable corresponding to the σ in the context). The intuition for the abstraction/closure case is that substituting the environment ‘undoes’ the effect of closure conversion on the context.

The compatibility relation is defined in Agda as follows:

```
data ~_~ : ∀ {Γ σ} → S.Lam σ Γ → T.Lam σ Γ → Set where

~V : ∀ {Γ σ} {x : Var σ Γ}
  → S.V x ~ T.V x

~A : ∀ {Γ σ τ} {L : S.Lam (σ ⇒ τ) Γ} {L† : T.Lam (σ ⇒ τ) Γ}
  {M : S.Lam σ Γ} {M† : T.Lam σ Γ}
  → L ~ L† → M ~ M†
  → S.A L M ~ T.A L† M†

~L : ∀ {Γ Δ σ τ} {N : S.Lam τ (σ :: Γ)}
  {N† : T.Lam τ (σ :: Δ)} {E : T.Subst Δ Γ}
  → N ~ T.subst (T.exts E) N†
  → S.L N ~ T.L N† E
```

We have defined the syntactic compatibility relation. The next section defines what it means for a relation to be a bisimulation.

5.1 Bisimulation

Bisimulation, as the name implies, is defined in terms on two simulations: one from source to target terms, and the other one from target to source terms.

In the following definitions speak about a two languages, A and B. Also, whenever simulations or bisimulations are mentioned, they are implicitly *lock-step*. The literature has example of more general simulations.

Definition. Given a relation \approx between terms of A and terms of B, we say that \approx is a **simulation** from A to B if and only if for all terms M and N in A, and M^\dagger in B, if M reduces in a single step to N, then there exists a term N^\dagger in B such that M^\dagger reduces to N^\dagger in a single step, and N is in the \approx relation with N^\dagger : $N \approx N^\dagger$.

The essence of simulation can be captured in a diagram.

$$\begin{array}{ccc} M & \longrightarrow & N \\ \approx \downarrow & & \downarrow \approx \\ M^\dagger & \longrightarrow & N^\dagger \end{array}$$

Recall that the *converse* of the relation \approx is a relation \approx' defined by $y \approx' x$ whenever $x \approx y$.

Definition. A relation \approx is a **bisimulation** if and only if it is a simulation and its converse is also a simulation.

In Agda, we instantiate the definition of simulation twice: once for a simulation from λ_{st} to λ_{cl} , and again for a simulation from λ_{cl} to λ_{st} :

`Rel = $\forall \{ \Gamma \ \sigma \} \rightarrow \text{S.Lam } \sigma \ \Gamma \rightarrow \text{T.Lam } \sigma \ \Gamma \rightarrow \text{Set}$`

`ST-Simulation : Rel \rightarrow Set`

`ST-Simulation $_ \approx _ = \forall \{ \Gamma \ \sigma \} \{ M \ N : \text{S.Lam } \sigma \ \Gamma \} \{ M^\dagger : \text{T.Lam } \sigma \ \Gamma \}$`
 `$\rightarrow M \approx M^\dagger \rightarrow M \text{ S.} \longrightarrow N$`

`$\rightarrow \exists [N^\dagger] ((N \approx N^\dagger) \times (M^\dagger \text{ T.} \longrightarrow N^\dagger))$`

`TS-Simulation : Rel \rightarrow Set`

`TS-Simulation $_ \approx _ = \forall \{ \Gamma \ \sigma \} \{ M : \text{S.Lam } \sigma \ \Gamma \} \{ M^\dagger \ N^\dagger : \text{T.Lam } \sigma \ \Gamma \}$`
 `$\rightarrow M \approx M^\dagger \rightarrow M^\dagger \text{ T.} \longrightarrow N^\dagger$`

`$\rightarrow \exists [N] ((N \approx N^\dagger) \times (M \text{ S.} \longrightarrow N))$`

Then we can provide an Agda definition of a bisimulation:

`Bisimulation : ($\forall \{ \Gamma \ \sigma \} \rightarrow \text{S.Lam } \sigma \ \Gamma \rightarrow \text{T.Lam } \sigma \ \Gamma \rightarrow \text{Set}$) \rightarrow Set`

`Bisimulation $_ \approx _ = \text{ST-Simulation } _ \approx _ \times \text{TS-Simulation } _ \approx _$`

To show that the compatibility relation is a bisimulation, we need to obtain lemmas about the interactions between the compatibility relation, values, renaming, and substitution.

5.2 Compatibility, values, renaming, and substitution

As discussed in [TODO], mechanising the meta-theory of a language involves proving lemmas about the interactions between various traversals and transformations, including renaming, substitution, and compilation phases. This is also the case for proving correctness with bisimulation, which requires establishing lemmas about the interplay between the compatibility relation, values, renaming, and substitution. In fact, proving those lemmas often constitutes the biggest effort in the entire proof. In Chapter 7, we reflect on the possibility of automating this effort with generic proving.

For each relevant property, we state it as an informal lemma, give its Agda statement, and its Agda proof.

Lemma. *Values commute with compatibility. If $M \sim M^\dagger$ and M is a value, then M^\dagger is also a value.*

The proof is by cases of term constructors.

```

~val : ∀ {Γ σ} {M : S.Lam σ Γ} {M† : T.Lam σ Γ}
  → M ~ M† → S.Value M
  → T.Value M†
~val ~V      ()
~val (~L ~N) S.V-L = T.V-L
~val (~A ~M ~N) ()

```

Lemma. *Renaming commutes with compatibility. If ρ is a renaming from Γ to Δ , and $M \sim M^\dagger$ are compatible terms in the context Γ , then the results of renaming M and M^\dagger with ρ are also compatible: $S.rename \rho M \sim T.rename \rho M^\dagger$.*

The proof is by induction on the similarity relation.

```

~rename : ∀ {Γ Δ σ} {M : S.Lam σ Γ} {M† : T.Lam σ Γ}
  → (ρ : Thinning Γ Δ) → M ~ M†
  → S.rename ρ M ~ T.rename ρ M†
~rename ρ ~V      = ~V
~rename ρ (~A ~M ~N) = ~A (~rename ρ ~M) (~rename ρ ~N)
~rename ρ (~L {N = N} ~N) with ~rename (T.ext ρ) ~N
... | ~ρN rewrite TT.lemma-~ren-L ρ E N† = ~L ~ρN

```

The variable and application cases are straightforward, but as ever, the abstraction case is more involved: it requires rewriting with an instantiation of the fusion lemma `rename∘subst`.

```

lemma-~ren-L : ∀ {Γ Δ Θ σ τ} (ρρ : Thinning Γ Θ) (ρσ : Subst Δ Γ) (N : Lam τ (σ :: Δ))
  → rename (ext ρρ) (subst (exts ρσ) N) ≡ subst (exts (rename ρρ <$> ρσ)) N

```

The final lemma is about the interplay between compatibility and substitution.

Definition. Suppose ρ and ρ^\dagger are two substitutions which take variables x in Γ to terms in Δ , such that for all x we have that $\text{lookup } \rho \ x \sim \text{lookup } \rho^\dagger \ x$. Then we say that ρ and ρ^\dagger are *pointwise compatible*.

Lemma. *Substitution commutes with compatibility. Suppose ρ and ρ^\dagger are two pointwise compatible substitutions. Then given compatible terms $M \sim M^\dagger$ in Γ , the results of applying ρ to M and ρ^\dagger to M^\dagger are also compatible: $S.\text{subst } \rho \ M \sim T.\text{subst } \rho^\dagger \ M^\dagger$.*

Pointwise similarity relation between substitutions ρ and ρ^\dagger is defined in Agda with $\sim\sigma$:

```
record _~σ_ {Γ Δ : Context} (ρ : S.Subst Γ Δ) (ρ† : T.Subst Γ Δ) : Set where
  field ρ~ρ† : ∀ → (x : Var σ Γ) → lookup ρ x ~ lookup ρ† x
```

We can show that pointwise similarity is preserved by applying exts to both substitutions:

```
~exts : ∀ {Γ Δ} {σ : Type} {ρ : S.Subst Γ Δ} {ρ† : T.Subst Γ Δ}
  → ρ ~σ ρ†
  → S.exts {τ = σ} ρ ~σ T.exts ρ†
ρ~ρ† (~exts ~ρ) z = ~V
ρ~ρ† (~exts {σ = σ} {ρ = ρ} ~ρ) (s x)
  = ~rename E.extend (ρ~ρ† ~ρ x)
```

In fact, extending pointwise-similar substitutions with similar terms preserves pointwise similarity:

```
_~•_ : ∀ {Γ Δ σ} {ρ : S.Subst Γ Δ} {ρ† : T.Subst Γ Δ}
  {M : S.Lam σ Δ} {M† : T.Lam σ Δ}
  → ρ ~σ ρ† → M ~ M†
  → ρ • M ~σ ρ† • M†
ρ~ρ† (ρ~σρ† ~• M~M†) z = M~M†
ρ~ρ† (ρ~σρ† ~• M~M†) (s x) = ρ~ρ† ρ~σρ† x
```

With the notion of pointwise similarity, we can prove that substitution commutes with similarity:

```
~subst : ∀ {Γ Δ τ} {ρ : S.Subst Γ Δ} {ρ† : T.Subst Γ Δ}
  {M : S.Lam τ Γ} {M† : T.Lam τ Γ}
  → ρ ~σ ρ† → M ~ M†
  → S.subst ρ M ~ T.subst ρ† M†
~subst ~ρ (~V {x = x}) = ρ~ρ† ~ρ x
~subst ~ρ (~A ~M ~N) = ~A (~subst ~ρ ~M) (~subst ~ρ ~N)
~subst {ρ† = ρ†} ~ρ (~L {N = N} ~N) with ~subst (~exts ~ρ) ~N
... | ~ρN rewrite TT.lemma-~subst-L ρ† E N† = ~L ~ρN
```

Just like in the lemma that renaming commutes with compatibility, the only non-trivial case is the one about abstractions/closures, which requires rewriting by an instantiation of the fusion lemma $\text{subst} \circ \text{subst}$.

$$\begin{aligned} \text{lemma-}\sim\text{subst-L} : & \forall \{ \Gamma \Delta \Theta \sigma \tau \} (\rho_1 : \text{Subst } \Gamma \Theta) (\rho_2 : \text{Subst } \Delta \Gamma) (N : \text{Lam } \tau (\sigma :: \Delta)) \\ & \rightarrow \text{subst } (\text{exts } \rho_1) (\text{subst } (\text{exts } \rho_2) N) \equiv \text{subst } (\text{exts } (\text{subst } \rho_1 <\$> \rho_2)) N \end{aligned}$$

With those three lemmas, showing that the compatibility relation is a bisimulation becomes straightforward.

5.3 Compatibility relation is a bisimulation

The proof that the compatibility relation \sim is a bisimulation consists of two proofs of simulations: $\text{st-sim} : \text{ST-Simulation } _ \sim _$ and $\text{ts-sim} : \text{TS-Simulation } _ \sim _$. Taken together, they prove that \sim is a bisimulation:

$$\begin{aligned} \text{bisim} : & \text{Bisimulation } _ \sim _ \\ \text{bisim} = & \text{st-sim} , \text{ts-sim} \end{aligned}$$

[TODO analyse the proof in detail or refer the reader to the technical appendix and to PLFA?]

5.4 Compatibility relation and closure conversion

We have showed that the compability relation is a bisimulation. But to link it to closure conversion and show that our implementation $_ \dagger$ is correct, we need to show that the graph relation of $_ \dagger$ is contained in the compatibility relation. In other words, if M is a source term, then we have $M \sim M \dagger$, where $M \dagger$ is the translation of M with $_ \dagger$. Thus, we need to show:

$$\begin{aligned} N \sim N \dagger : & \forall \{ \Gamma A \} (N : \text{S.Lam } A \Gamma) \\ & \rightarrow N \sim N \dagger \end{aligned}$$

The proof of this is too long to discuss here, but the reader can find it in the technical appendix of this report.

However, as we argued, the graph relation of any well-behaved closure conversion function is contained within the compatibility relation. It is instructive to see that this is the case for a simpler version of closure conversion, namely, a trivial closure conversion algorithm simple-cc , which uses full contexts as environments (through identity substitutions id-subst):

$$\begin{aligned} \text{simple-cc} : & \forall \{ \Gamma \sigma \} \rightarrow \text{S.Lam } \sigma \Gamma \rightarrow \text{T.Lam } \sigma \Gamma \\ \text{simple-cc } (\text{S.V } x) = & \text{T.V } x \end{aligned}$$

$\text{simple-cc } (\mathbf{S.A} \ M \ N) = \mathbf{T.A} \ (\text{simple-cc } M) \ (\text{simple-cc } N)$
 $\text{simple-cc } (\mathbf{S.L} \ N) = \mathbf{T.L} \ (\text{simple-cc } N) \ \mathbf{T.id-subst}$

Indeed, the graph relation of `simple-cc` is contained in the similarity relation. The proof is by straightforward induction; in the abstraction case, we need to argue that applying an identity substitution leaves the argument term unchanged.

$\text{simple-cc} \rightarrow \text{sim} : \forall \{\Gamma \ \sigma\} \ (N : \mathbf{S.Lam} \ \sigma \ \Gamma) \rightarrow N \sim \text{simple-cc } N$
 $\text{simple-cc} \rightarrow \text{sim} \ (\mathbf{S.V} \ x) = \sim \mathbf{V}$
 $\text{simple-cc} \rightarrow \text{sim} \ (\mathbf{S.A} \ f \ e) = \sim \mathbf{A} \ (\text{simple-cc} \rightarrow \text{sim} \ f) \ (\text{simple-cc} \rightarrow \text{sim} \ e)$
 $\text{simple-cc} \rightarrow \text{sim} \ (\mathbf{S.L} \ b) = \sim \mathbf{L} \ g$
 where
 $h : \forall \{\Gamma \ \sigma \ \tau\} \ (M : \mathbf{T.Lam} \ \sigma \ (\tau :: \Gamma)) \rightarrow \mathbf{T.subst} \ (\mathbf{T.exts} \ \mathbf{T.id-subst}) \ M \equiv M$
 $h \ M =$
 begin
 $\mathbf{T.subst} \ (\mathbf{T.exts} \ \mathbf{T.id-subst}) \ M$
 $\equiv \langle \text{cong } (\lambda e \rightarrow \mathbf{T.subst} \ e \ M) \ (\text{sym } (\text{env-extensionality} \ \mathbf{TT.exts-id-subst})) \rangle$
 $\mathbf{T.subst} \ \mathbf{T.id-subst} \ M$
 $\equiv \langle \mathbf{TT.subst-id-id} \ M \rangle$
 M
 ■
 $g : b \sim \mathbf{T.subst} \ (\mathbf{T.exts} \ \mathbf{T.id-subst}) \ (\text{simple-cc } b)$
 $g \ \text{rewrite } h \ (\text{simple-cc } b) = \text{simple-cc} \rightarrow \text{sim} \ b$

TBC

Chapter 6

Proof by logical relations

As we mentioned in ??, there are two standard methods for proving operational correctness of a translation: bisimulations and logical relations [TODO wording from Dreyer’s paper]. Chapter 4 discussed an Agda mechanisation of a proof of bisimulation for [TODO wording] closure conversion. This chapter presents a mechanisation of the other proof method, [by/with?] logical relations?

The chapter starts with a presentation of an modified formalisation of the source and target languages of closure conversion. Then, a pen-and-paper proof by logical relations is given, and finally, its Agda formalisation.

6.1 Alternative formalisation of the intermediate languages

This section presents an alternative formalisation of the source and target languages of closure conversion. We call the new formalisation of the source language λ_{st}' , and the new formalisation of the target language - λ_{cl}' . Compared with λ_{st} and λ_{cl} in Chapter 4, λ_{st}' and λ_{cl}' are different in two ways.. Firstly, the distinction between values and non-values is made explicit in the definition of terms in λ_{st}' and λ_{cl}' , replacing a predicate on terms in λ_{st} and λ_{cl} . Secondly, we give big-step semantics for λ_{st}' and λ_{cl}' , in contrast to small-step semantics for λ_{st} and λ_{cl} . These two differences simplify mechanisation of a proof by logical relation.

These improvements in formalisation are inspired by an Agda formalisation accompanying [6].

[TODO maybe only discuss STLC?]

The definitions of types, contexts, variables as proofs of context membership, and environments, are the same as for λ_{st} and λ_{cl} in the previous chapter. The definition of language expressions is different, however, in that it makes an explicit distinction between values Val and non-values Trm . This is achieved by indexing the Exp data type by a $Kind$:


```

data Kind : Set where
  'val 'trm : Kind

data Exp : Kind → Type → Context → Set

Trm : Type → Context → Set
Trm = Exp 'trm

Val : Type → Context → Set
Val = Exp 'val

infixl 5 _'$_

data Exp where

  -- values
  'var : ∀ {Γ σ} → Var σ Γ → Val σ Γ
  'λ : ∀ {Γ σ τ} → Trm τ (σ :: Γ) → Val (σ ⇒ τ) Γ

  -- non-values (a.k.a. terms)
  _'$_ : ∀ {Γ σ τ} → Val (σ ⇒ τ) Γ → Val σ Γ → Trm τ Γ
  'let : ∀ {Γ σ τ} → Trm σ Γ → Trm τ (σ :: Γ) → Trm τ Γ
  'val : ∀ {Γ σ} → Val σ Γ → Trm σ Γ

```

Notice that there are two new constructors for language expressions. The first one is 'val, which takes a value Val to a term Trm and thus makes it possible to use values in positions where terms are expected. The second is 'let, which is a standard let construct. The let is necessary to make the evaluation order explicit: function application applies a value to a value, so nested computations need to be factored out and bound as values by a let expression. This representation is known as A-normal form [8], and it is used for λ_{st} and λ_{cl} as it simplifies the definition as big-step semantics.

Definition of renaming and substitution are similar to those for λ_{st} , so we do not include the updated versions here. Instead, we define aliases for closed values Val_0 and closed terms Trm_0 (typed in an empty context):

```

Exp0 : Kind → Type → Set
Exp0 k τ = Exp k τ []

Trm0 : Type → Set
Trm0 = Exp0 'trm

Val0 : Type → Set
Val0 = Exp0 'val

```

Like it was mentioned, the semantics of λ_{st} are defined as big-step semantics. Given a term M and a value V, the inductive definition $M \Downarrow V$ states the conditions for M to reduce to a value V:

```

data  $\rightarrow_1$  :  $\forall \rightarrow \text{Trm}_0 \sigma \rightarrow \text{Trm}_0 \sigma \rightarrow \text{Set where}$ 
 $\rightarrow_1 \text{app}$  :  $\forall \{\sigma \tau\} \{M : \text{Trm } \tau (\sigma :: [])\} \{V : \text{Val}_0 \sigma\} \rightarrow \lambda M \$ V \rightarrow_1 M [V]$ 

data  $\Downarrow$  :  $\forall \rightarrow \text{Trm}_0 \sigma \rightarrow \text{Val}_0 \sigma \rightarrow \text{Set where}$ 
 $\Downarrow \text{val}$  :  $\forall \{V : \text{Val}_0 \sigma\} \rightarrow \text{val } V \Downarrow V$ 
 $\Downarrow \text{app}$  :  $\forall \{\sigma \tau\} \{M : \text{Trm } \tau (\sigma :: [])\} \{V : \text{Val}_0 \sigma\} \{U : \text{Val}_0 \tau\}$ 
 $\rightarrow M [V] \Downarrow U \rightarrow \lambda M \$ V \Downarrow U$ 
 $\Downarrow \text{let}$  :  $\forall \{\sigma \tau\} \{M : \text{Trm}_0 \sigma\} \{N : \text{Trm } \tau (\sigma :: [])\} \{U : \text{Val}_0 \sigma\} \{V : \text{Val}_0 \tau\}$ 
 $\rightarrow M \Downarrow U \rightarrow N [U] \Downarrow V \rightarrow \text{let } M N \Downarrow V$ 
 $\Downarrow \text{step}$  :  $\forall \{M M' : \text{Trm}_0 \sigma\} \{V : \text{Val}_0 \sigma\} \rightarrow M \rightarrow_1 M' \rightarrow M' \Downarrow V \rightarrow M \Downarrow V$ 

```

It is worth explaining the $\Downarrow \text{step}$ constructor and the $M \rightarrow_1 M'$ data type. The $M \rightarrow_1 M'$ data type describes the small-step reduction relation and has a single constructor which captures beta reduction for functions. The $\Downarrow \text{step}$ constructor is similar to the transitive closure of the small-step reduction relation: if M reduces to M' in a single step, and M' reduces to V in multiple steps, then M reduces to V in multiple steps. [TODO maybe complete the explanation].

Finally, ... [TODO what to make of a non-terminating proof of termination?]

```

{-# TERMINATING #-}
sn :  $\forall (N : \text{Trm}_0 \sigma) \rightarrow \Sigma [V \in \text{Val}_0 \sigma] (N \Downarrow V)$ 
sn ('var () '$ _)
sn (' $\lambda M \$ V$ ) with sn (M [V])
sn (' $\lambda M \$ V$ ) | U , M[V] $\Downarrow$ U = U ,  $\Downarrow \text{step} \rightarrow_1 \text{app } M[V] \Downarrow U$ 
sn ('let M N) with sn M
sn ('let M N) | U , M $\Downarrow$ U with sn (N [U])
sn ('let M N) | U , M $\Downarrow$ U | V , N $\Downarrow$ V = V ,  $\Downarrow \text{let } M \Downarrow U N \Downarrow V$ 
sn ('val V) = V ,  $\Downarrow \text{val}$ 

```

Differences between λcl and $\lambda \text{cl}'$ are analogous.

6.2 Proving correctness of a translation with logical relations

This section defines two relations between terms of $\lambda \text{st}'$ and $\lambda \text{cl}'$, one syntactic, and one operational. The syntactic relation, which we call a *compatibility relation*, subsumes the graph relation of any closure conversion, and the operational relation captures the intuition that similar terms reduce to similar values. We prove that the syntactic relation entails the operational relation, and thus, that any well-behaved closure conversion function preserves operational correctness. The proof uses type-indexed logical relations and is inspired by a sketch of a similar proof from [7].

For brevity, we do not duplicate translation functions from λst to $\lambda \text{st}'$, but just like we did in Section ??, we expect that every the graph relation of every well-behaved translation from $\lambda \text{st}'$ to $\lambda \text{cl}'$ will be contained in the compatibility relation \cong . In general,

given a term M_1 in λ_{st}' and a term M_2 in λ_{cl}' , M_1 and M_2 are in the compatibility relation ($M_1 \cong M_2$) when their subterms are in the compatibility relation. In the special case of abstractions/closures, the closure body is renamed with the environment in the premise.

```

infix 4 _≅_
data _≅_ : ∀ {Γ σ k} → S.Exp k σ Γ → T.Exp k σ Γ → Set where

-- values

~var : ∀ {Γ σ} {x : Var σ Γ}
-----
→ S.'var x ≅ T.'var x

~λ : ∀ {Γ Δ σ τ} {N1 : S.Trm τ (σ :: Γ)} {N2 : T.Trm τ (σ :: Δ)} {E : T.Subst Δ Γ}
→ N1 ≅ T.subst (T.exts E) N2
-----
→ S.'λ N1 ≅ T.'λ N2 E

-- terms

_~$ : ∀ {Γ σ τ} {L : S.Val (σ ⇒ τ) Γ} {L† : T.Val (σ ⇒ τ) Γ}
      {M : S.Val σ Γ} {M† : T.Val σ Γ}
→ L ≅ L†
→ M ≅ M†
-----
→ L S.$ M ≅ L† T.$ M†

~let : ∀ {Γ σ τ} {M1 : S.Trm σ Γ} {M2 : T.Trm σ Γ}
      {N1 : S.Trm τ (σ :: Γ)} {N2 : T.Trm τ (σ :: Γ)}
→ M1 ≅ M2
→ N1 ≅ N2
-----
→ S.'let M1 N1 ≅ T.'let M2 N2

~val : ∀ {Γ σ} {M1 : S.Val σ Γ} {M2 : T.Val σ Γ}
→ M1 ≅ M2
-----
→ S.'val M1 ≅ T.'val M2

```

The compatibility relation provides handy inductive hypotheses in the proof by logical relations.

While the compatibility relation captures syntactic correspondence, we need another relation on (closed) language expressions which captures operational correspondence. In fact, this relation \Leftrightarrow is comprised of (1) a relation \sim between closed source terms and closed target terms, and (2) a relation \approx between closed source values and closed

target values. Those two relations are defined by mutual induction and induction on types [TODO both at the same time?] as follows:

$\tau \ni M_1 \sim M_2$ iff $M_1 \Downarrow V_1, M_2 \Downarrow V_2$, and $\tau \ni V_1 \approx V_2$

$\sigma \Rightarrow \tau \ni U_1 \approx U_2$ iff for all $\sigma \ni V_1 \approx V_2$, we have $\tau \ni U_1 \Downarrow V_2 \sim U_2 \Downarrow V_2$

There is no case for \approx at the ground type α as only variables can have the ground type, and the values in \approx are closed.

In Agda, \sim and \approx are defined as specialisations of a relation \Leftrightarrow on expressions of λ st and λ cl.

```
{-# NO_POSITIVITY_CHECK #-}
data _⇔_ : {k : Kind} {τ : Type} → S.Exp₀ k τ → T.Exp₀ k τ → Set

_~_ : ∀ → S.Trm₀ τ → T.Trm₀ τ → Set
_~_ = _⇔_

_≈_ : ∀ → S.Val₀ τ → T.Val₀ τ → Set
_≈_ = _⇔_

data _⇔_ where

  -- values

  ≈λ : ∀ {Δ σ τ} {M₁ : S.Trm τ (σ :: [])}
        {M₂ : T.Trm τ (σ :: Δ)} {E : T.Subst Δ []}
        → ({V₁ : S.Val₀ σ} {V₂ : T.Val₀ σ}
            → V₁ ≈ V₂ → M₁ [ V₁ ] ~ T.subst (E • V₂) M₂)
        → S.λ M₁ ≈ T.λ M₂ E

  -- terms

  ~Trm : ∀ {N₁ : S.Trm₀ σ} {N₂ : T.Trm₀ σ}
          {V₁ : S.Val₀ σ} {V₂ : T.Val₀ σ}
          → N₁ S.↓ V₁
          → N₂ T.↓ V₂
          → V₁ ≈ V₂
          → N₁ ~ N₂
```

Finally, we extend the \approx relation to source and target substitution environments, similar to what we did in Section ??:

```
record _•≈_ {Γ : List Type}
  (ρs : S.Subst Γ []) (ρt : T.Subst Γ []) : Set where
  constructor packR
```

$$\begin{aligned} \text{field lookup}^R &: \{\sigma : \text{Type}\} (v : \text{Var } \sigma \Gamma) \\ &\rightarrow \text{lookup } \rho^s v \approx \text{lookup } \rho^t v \end{aligned}$$

We also provide a function \bullet^R which extends two related substitution environments with a related pair of values:

$$\begin{aligned} _ \bullet^R _ &: \forall \{\Gamma \tau\} \\ &\quad \{\rho^s : \text{S.Subst } \Gamma []\} \{\rho^t : \text{T.Subst } \Gamma []\} \\ &\quad \{N_1 : \text{S.Val}_0 \tau\} \{N_2 : \text{T.Val}_0 \tau\} \\ &\rightarrow \rho^s \bullet \approx \rho^t \\ &\rightarrow N_1 \approx N_2 \\ &\quad \text{-----} \\ &\rightarrow \rho^s \bullet N_1 \bullet \approx \rho^t \bullet N_2 \\ \text{lookup}^R (\rho^R \bullet^R \approx N) \text{z} &= \approx N \\ \text{lookup}^R (\rho^R \bullet^R \approx N) (\text{s } x) &= \text{lookup}^R \rho^R x \end{aligned}$$

With those definitions at hands, we can state the Fundamental Theorem of Logical Relations. Recall that we are trying to show that, given a closed term M_1 in $\lambda\text{st}'$ and a closed term M_2 in λcl , if $M_1 \cong M_2$, then $M_1 \sim M_2$ (*). The Fundamental Theorem is:

$$\begin{aligned} \text{fund} &: \forall \{\Gamma \sigma k\} \{M_1 : \text{S.Exp } k \sigma \Gamma\} \{M_2 : \text{T.Exp } k \sigma \Gamma\} \\ &\quad \{\rho^s : \text{S.Subst } \Gamma []\} \{\rho^t : \text{T.Subst } \Gamma []\} \\ &\rightarrow \rho^s \bullet \approx \rho^t \\ &\rightarrow M_1 \cong M_2 \\ &\quad \text{-----} \\ &\rightarrow \text{S.subst } \rho^s M_1 \Leftrightarrow \text{T.subst } \rho^t M_2 \end{aligned}$$

Thus, the Fundamental Theorem is a stronger statement than (*), and one which instantiated with closed terms and empty substitution environments yields precisely (*).

We do not include the Agda proof here as it is not very readable; instead, we present several cases on paper: TODO

Chapter 7

Reflections and evaluation

This project is a case study on verification of transformations of functional programs using two different techniques: bisimulations and logical relations. The implemented transformation is closure conversion. Both proofs of operational correctness are mechanised with state-of-the-art techniques.

The **original objectives** of the project were:

1. To implement a compiler transformation for a variant of simply-typed lambda calculus in Agda.
2. To use scope-safe and well-typed representation for the object languages.
3. To prove that the transformation is correct: that the output program of the transformation behaves ‘the same’ as the input program.
4. To use generic programming techniques from ACMM.

The **contributions** of this project are as follows:

1. All the original objectives were achieved.
2. It is demonstrated that languages with closures and closure conversion are problematic for current state-of-the-art techniques for mechanising language meta-theory.

(Objective 1) Capturing the essence of closure conversion The implemented transformation — closure conversion — requires a different source and target language. While the formalisation of the source language is largely borrowed from ACMM, and the formalisation of the target language is similar except for the difference between abstractions and closures, this project’s contribution was to capture the essence of closure conversion in what we believe is the simplest and most elegant way possible.

Inherently typed closures A traditional representation of closure conversion replaces variables in the source program with references to a record containing the environment in the target program. This project’s use of scope-safe and well-typed terms

allowed for a more elegant solution where the closure body is typed in a context corresponding to the closure's environment, and variables remain variables.

Closure environments as substitution environments Furthermore, while a closure environment is traditionally represented as a record which stores environment values, this project captures the essence of an environment by representing it as a substitution environment, i.e. a mapping from variables to values.

Existential types for closure environments As this report points out, closure environment must have existential types in order for a program with closures to be well-typed. This observation was made by [7], which deals with this fact by equipping the closure language with existential types. This project uses a different, arguably simpler approach, whereby closure environment are existentially typed *in the meta language* (*Agda*), which allows us to keep object language types simple.

Comparison with traditional closure conversion In comparison with traditional closure conversion which represents environments as records, this formulation, which represents closure environments as substitution environments, i.e. meta-language functions, is further removed from the eventual target, which is machine code. But one can imagine a subsequent compilation phase which replaces substitution environments with records, and variables with record lookups (the object language would need existential types then). In general, splitting the compilation process into many specialised passes facilitates verification, as each compilation phase is easier to verify, and composing correctness results about phases gives rise to a end-to-end correctness result.

(Objective 2) Scope-safe and well-typed representation Both the source and target language have scope-safe and well-typed representation, which were possible thanks to using dependently-typed *Agda* as the meta language. Using inherently scoped and typed terms has many benefits, which include the fact that when programs are synonymous with their typing derivations, transformations on programs are synonymous with proofs of type preservation. Additionally, many techniques for reasoning about operational correctness are type directed, e.g. the type-indexed logical relations which we used, and inherently typed representations are well-matched to such techniques.

(Objective 3) The closure conversion preserves operational correctness This project uses two standard techniques to show that the implemented closure conversion is correct: bisimulation and logical relations. In an informal setting of pen-and-paper proofs, both of those techniques have rather straightforward proofs. However, mechanisation of those proofs involves proving several lemmas about the interactions between renaming, substitution, closure conversion, and the compatibility relation.

Mechanising the meta-theory of a language As observed in ACMM, mechanising the meta-theory of a language most often requires proving lemmas about the interactions between different transformations, or semantics, like renaming and substitutions. ACMM singles out synchronisation lemmas, which relate two semantics (e.g. for every renaming there exists a substitution which behaves the same), and fusion lemmas (e.g. for every composition of two substitutions, there exists a substitution which behaves the same).

(Objective 4) ACMM ACMM exploit similarity between various traversals (semantics) on simply typed lambda calculus (STLC) to come up with a generic way to prove synchronisation and fusion lemmas for STLC. Indeed, this project uses fusion lemmas about STLC from ACMM in the proof with logical relations.

Intermediate languages other than STLC require their own definitions of renaming and substitution, and proofs of correctness lemmas. For example, the proofs of operational correctness with bisimulations and logical relations depend on four fusion lemmas relating renaming and substitution for the language with closures. In fact, proving those correctness lemmas was the biggest effort in the whole proof.

Possible remedy: AACMM and generic programming The problem of having to define renaming and substitution for each new language, and proving correctness lemmas about the interactions between renaming, substitution, and transformations, is addressed by a follow-up paper, which we will refer to as AACMM [2]. AACMM provides a way to supply a definition of a syntax with bindings, and then derives meta-theoretical correctness lemmas from that definition. The paper repository contains an example of an elaboration whose source is a language with a `let` construct, and whose target is simply-typed lambda calculus with `let`-expressions inlined. The example demonstrates how a proof of simulation is drastically simplified thanks to the use of the AACMM library and its generic programming capabilities.

Feasibility of closure conversion in AACMM AACMM demonstrates that transformations like `let`-inlining and CPS conversion can be expressed in their generic framework. They pose an open question about which compilation passes can be implemented generically. Unfortunately, this work suggests that closure conversion might not fit well into the AACMM framework. Specifically, the closure language in this project — with aforementioned features like syntax being mutually dependent on substitution environments, or environments being existentially quantified in the meta language (Agda) — is not expressible as an AACMM generic syntax. The traditional representation of a language with closures — with environments as records and existential types in the object language — would not fit either as syntaxes in AACMM cannot contain existential types.

Bisimulation vs logical relations [TODO what could I say about the pros and cons of both?]

Summary This work and ACMM/AACMM are both concerned with mechanising the meta-theory of languages, and applying this metatheory to reason about the language. While AACMM shows that a certain class of languages / syntaxes can be treated generically, this work contains a negative result which indicates that a language with closures might not benefit from current techniques for relieving the burden of mechanising meta-theory. This is an open question, however, whether there exist feasible generic syntaxes which would encompass a language with closures, or whether an alternative formalisation of a language with closures exists which is compatible with AACMM.

[TODO the conclusions of this chapter depend on my understanding of AACMM — is it correct?]

Chapter 8

Relationship to the UG4 project

This work is a natural continuation of the UG4 project, but it admittedly takes the project in a new direction. In terms of its goals, the UG4 project was concerned with program derivations. Derivations are distinct from program transformations or traversals, as found in compilers or in this UG5 project.

Program transformations like closure conversion, or continuation passing style (CPS) transformation, have two distinct characteristics.

8.1 Program transformation vs derivation

Firstly, the source and target languages can be different, e.g. the source language may have lambda abstractions with free variables, and the target language may have closures with environments. Of course, many transformations in compilers happen within the same language, e.g. constant expression folding.

Secondly, compiler transformations are usually characterised by replacing one program construct with another, in a one-way fashion, e.g. lambda abstractions with closures. It would be strange if changes happened both ways. One might imagine a language with both lambda abstractions and closures, and a transformation which replaces some closures with lambda abstractions, and some abstraction with closures, according to arbitrary rules. It is difficult to see how such a transformation would be useful in a compiler.

Of course, one might imagine a transformation which replaces some occurrences of constructs A and constructs B, and vice versa, in order to optimise the program. But any such optimising transformation is still guided by some measure of performance.

In contrast, program derivation consists of transforming a program in arbitrary places, and using arbitrary rules, with a specific goal of obtaining one program from another, so that the obtained program has some desirable features, like efficiency, or even faster asymptotic running time. Importantly, the derivation happens within a single language.

(TODO but specification can be in terms of relations, which are not part of the language).

8.2 Program derivations in the UG4 project

The UG4 project analyses two instances of program derivation in detail. The first one is a derivation of an efficient implementation of the maximum segment sum problem (MSS). While the input specification (which is also a runnable program) runs in cubic time in the length of the input list, the output program runs in linear time. The asymptotic speed-up is achieved by applying several "rewrite rules" involving higher-order functions on lists such as `map`, `foldr`, and `filter`.

The second case study involved a derivation for a program for matrix-vector multiplication. The input program takes a dense matrix, and the output program takes a sparse matrix in the compressed sparse row (CSR) format. Or, to be precise, the input program which acts on a dense matrix, is transformed into a composition of two programs: (a) a conversion from a dense to a CSR-sparse matrix, and (b) a matrix-vector multiplication program which acts on a CSR-sparse matrix. This is because, as a rule, the input and output types of the program must stay the same in the course of the derivation. This second derivation was similarly accomplished with "rewrite rules" involving higher-order functions.

8.3 Rewrite rules

The notion of a rewrite rule is central to program derivation. A classical example of a rewrite rule for a function program is one stating: "a composition of mappings is a mapping with a composition":

```
forall f g xs. map f (map g xs) == map (f . g) xs (*)
```

where `f`, `g`, `xs` are metavariables. An application of a rewrite rule consists of unifying the LHS of the rule with a subterm of the program (and thus obtaining a substitution σ), and then replacing the subterm with the RHS of the rule, instantiated with the substitution σ .

Notably, such simple form of a rewrite rule only supports first-order abstract syntax trees, but not higher-order abstract syntax. (TODO elaborate on what it would mean to have a context and go under a binder in a rewrite rule). (TODO can't express conditions)

In their simplest form, a program derivation is a sequence of intermediate forms of the program, intertwined with rewrite rules which justify each step of the derivation. For examples, the reader may consult the UG4 project [?].

8.4 Program derivations in compilers and their limitations

While we argued that compiler transformations and program derivations are distinct in their character, the lines may arguably be blurry at times. A good example of this is the support for rewrite rules in GHC, a Haskell compiler. A Haskell programmer may specify a rule like $(*)$ as part of the code, and in one of early compilation passes, GHC will apply the rule wherever possible (i.e. replace the occurrences of the LHS with the RHS). Such compiler pass may be considered an instance of program derivation, and it would go some way towards deriving the aforementioned efficient implementation for the maximum segment sum problem (MSS).

However, rewriting as implemented in compilers is too limited to carry out most derivations. To see one limitation, consider a derivation which needs to apply the rule $(*)$ right to left: replace an occurrence of the RHS with the LHS. But clearly, unguided application of rewrite rules can only be one way, otherwise it would not terminate.

Another limitation is the fact that "the right" derivation can require that rules be applied in a specific order. Thus, rewriting a program becomes a search problem, where rewrites are applied until a program satisfying some (performance) objective is found. This is the approach taken by the Lift compiler [?]

Yet another limitation is that many derivations elude the notion of an objective function, thus rendering a search futile. It seems that human insight is required to guide a derivation.

A final limitation is that there are conceivable rewrite rules which are only applicable when certain conditions are met. Such conditions could range algebraic properties of operators to general predicates and relations on terms. And these are undecidable in general (TODO how to phrase this).

These considerations, taken together, mean that the technique of program derivation is more useful to the programmer than the compiler. Benefits of employing a sort of program derivation (more or less formal) for the programmer include: (a) a structured process of obtaining an implementation from specification, (b) greater confidence about correctness of an implementation, (c) possibility of discovering further optimisations, and finally (d) a framework for a proof of correctness. This last use case could be explained as follows: suppose we can prove the correctness of the "specification" program, and the correctness of each rewrite rule. Then we can obtain correctness of the "implementation" program.

8.5 Implementation of rewriting in the UG4 project

The UG4 project included a purpose-built framework for specifying derivations. The framework included:

1. A simple functional language with parametricity. The language is point-free, that is, based on function composition rather than on lambda abstractions. This was because variables and abstraction are difficult to implement correctly, as demonstrated by this UG5 project, and even more difficult to rewrite.
2. A type-checker for the language.
3. Rewriting functionality and declaring derivations as sequences of rewrites.
4. An interpreter for the language, which was used to empirically verify claims about performance gains from derivations.

Writing the framework was a good exercise in implementing routine parts of compiler front-ends, such as type checking and unification. Writing it in Scala made sense given the stretch objective of compiling the language to Lift, which was not realised, however.

Importantly, rewrite rules were stated without justification, much as postulates in Agda. One could prove the rules externally – but then one is pressed to ask, why not express a derivation in a proof assistant, which supports unification and rewriting natively? Indeed, with hindsight, we can say with certainty that a proof assistant is perfectly suited for the job, its only downside being that it requires considerable expertise, which I did not have during my fourth year. (TODO I/me?)

8.6 Program derivation in Agda: sparse matrix-vector multiplication

To complete last year's work, we conduct a derivation of the program for matrix-vector multiplication which acts on CSR-sparse matrices. Unlike last year, we can now provide proofs of individual rewrite rules. Indeed, some proofs are quite involved. TODO whether and how to do it.

Bibliography

- [1] Andreas Abel, Brigitte Pientka, David Thibodeau, and Anton Setzer. Copatterns: programming infinite structures by observations. In Roberto Giacobazzi and Radhia Cousot, editors, *The 40th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, POPL '13, Rome, Italy - January 23 - 25, 2013*, pages 27–38. ACM, 2013.
- [2] Guillaume Allais, Robert Atkey, James Chapman, Conor McBride, and James McKinna. A type and scope safe universe of syntaxes with binding: their semantics and proofs. *PACMPL*, 2(ICFP):90:1–90:30, 2018.
- [3] Guillaume Allais, James Chapman, Conor McBride, and James McKinna. Type-and-scope safe programs and their proofs. In Yves Bertot and Viktor Vafeiadis, editors, *Proceedings of the 6th ACM SIGPLAN Conference on Certified Programs and Proofs, CPP 2017, Paris, France, January 16-17, 2017*, pages 195–207. ACM, 2017.
- [4] Thorsten Altenkirch and Bernhard Reus. Monadic presentations of lambda terms using generalized inductive types. In Jörg Flum and Mario Rodríguez-Artalejo, editors, *Computer Science Logic, 13th International Workshop, CSL '99, 8th Annual Conference of the EACSL, Madrid, Spain, September 20-25, 1999, Proceedings*, volume 1683 of *Lecture Notes in Computer Science*, pages 453–468. Springer, 1999.
- [5] Conor McBride. Type-preserving renaming and substitution. 2005.
- [6] Craig McLaughlin, James McKinna, and Ian Stark. Triangulating context lemmas. In June Andronick and Amy P. Felty, editors, *Proceedings of the 7th ACM SIGPLAN International Conference on Certified Programs and Proofs, CPP 2018, Los Angeles, CA, USA, January 8-9, 2018*, pages 102–114. ACM, 2018.
- [7] Yasuhiko Minamide, J. Gregory Morrisett, and Robert Harper. Typed closure conversion. In Hans-Juergen Boehm and Guy L. Steele Jr., editors, *Conference Record of POPL'96: The 23rd ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages, Papers Presented at the Symposium, St. Petersburg Beach, Florida, USA, January 21-24, 1996*, pages 271–283. ACM Press, 1996.
- [8] Amr Sabry and Matthias Felleisen. Reasoning about programs in continuation-passing style. In *LISP and Functional Programming*, pages 288–298, 1992.

- [9] Philip Wadler. Programming language foundations in agda. In Tiago Massoni and Mohammad Reza Mousavi, editors, *Formal Methods: Foundations and Applications - 21st Brazilian Symposium, SBMF 2018, Salvador, Brazil, November 26-30, 2018, Proceedings*, volume 11254 of *Lecture Notes in Computer Science*, pages 56–73. Springer, 2018.

Chapter 9

Technical appendix

9.1 Minimising closure conversion and the similarity relation

Below is the Agda proof that the graph relation of the minimising closure conversion function is contained in the similarity relation.

```
_† : ∀ {Γ A} → S.Lam A Γ → T.Lam A Γ
_† M with cc M
_† M | ∃[ Δ ] Δ ⊆ Γ ∧ N = T.rename (⊆→p Δ ⊆ Γ) N
```

```
helper-2 : ∀ {Γ A} (x : Var A Γ)
  → lookup (⊆→p (Var→⊆ x)) z ≡ x
helper-2 z = refl
helper-2 (s x) = cong s (helper-2 x)
```

```
helper-3 : ∀ {Δ1 Γ1 Γ} (Δ1 ⊆ Γ1 : Δ1 ⊆ Γ1) (Γ1 ⊆ Γ : Γ1 ⊆ Γ)
  → select (⊆→p Δ1 ⊆ Γ1) (⊆→p Γ1 ⊆ Γ) ≡E ⊆→p (⊆-trans Δ1 ⊆ Γ1 Γ1 ⊆ Γ)
eq (helper-3 base base) ()
eq (helper-3 Δ1 ⊆ Γ1 (skip Γ1 ⊆ Γ)) x
  = cong s (eq (helper-3 Δ1 ⊆ Γ1 Γ1 ⊆ Γ) x)
eq (helper-3 (skip Δ1 ⊆ Γ1) (keep Γ1 ⊆ Γ)) x
  = cong s (eq (helper-3 Δ1 ⊆ Γ1 Γ1 ⊆ Γ) x)
eq (helper-3 (keep Δ1 ⊆ Γ1) (keep Γ1 ⊆ Γ)) z
  = refl
eq (helper-3 (keep Δ1 ⊆ Γ1) (keep Γ1 ⊆ Γ)) (s x)
  = cong s (eq (helper-3 Δ1 ⊆ Γ1 Γ1 ⊆ Γ) x)
```

```
helper-4 : ∀ {Δ1 Γ1 Γ τ}
  (Δ1 ⊆ Γ1 : Δ1 ⊆ Γ1) (Γ1 ⊆ Γ : Γ1 ⊆ Γ)
  (Δ1 ⊆ Γ : Δ1 ⊆ Γ) (M† : T.Lam τ Δ1)
  → ⊆-trans Δ1 ⊆ Γ1 Γ1 ⊆ Γ ≡ Δ1 ⊆ Γ
```



```

→ T.rename (⊆→p Γ1⊆Γ) (T.rename (⊆→p Δ1⊆Γ1) M†)
≡ T.rename (⊆→p Δ1⊆Γ) M†
helper-4 Δ1⊆Γ1 Γ1⊆Γ Δ1⊆Γ M† well =
begin
  T.rename (⊆→p Γ1⊆Γ) (T.rename (⊆→p Δ1⊆Γ1) M†)
≡ ⟨ rename◦rename (⊆→p Δ1⊆Γ1) (⊆→p Γ1⊆Γ) M† ⟩
  T.rename (select (⊆→p Δ1⊆Γ1) (⊆→p Γ1⊆Γ)) M†
≡ ⟨ cong (λ e → T.rename e M†)
      (env-extensionality (helper-3 Δ1⊆Γ1 Γ1⊆Γ)) ⟩
  T.rename (⊆→p (⊆-trans Δ1⊆Γ1 Γ1⊆Γ)) M†
≡ ⟨ cong (λ e → T.rename (⊆→p e) M†) well ⟩
  T.rename (⊆→p Δ1⊆Γ) M†

```

```

{-# TERMINATING #-}
helper-5 : ∀ {Γ Δ σ τ} (Δ⊆Γ : Δ ⊆ Γ) (N : T.Lam σ (τ :: Δ))
→ T.subst (T.exts (T.rename (⊆→p Δ⊆Γ) <$> T.id-subst)) N
≡ T.rename (⊆→p (keep Δ⊆Γ)) N
helper-5 Δ⊆Γ (T.V x) with x
helper-5 Δ⊆Γ (T.V x) | z = refl
helper-5 Δ⊆Γ (T.V x) | s x' = refl
helper-5 Δ⊆Γ (T.A M N)
= cong2 T.A (helper-5 Δ⊆Γ M) (helper-5 Δ⊆Γ N)
helper-5 Δ⊆Γ (T.L N E) = cong (T.L N) h
where
h : T.subst (T.exts (T.rename (⊆→p Δ⊆Γ) <$> T.id-subst)) <$> E
≡ _<$>_ {W = T.Lam} (T.rename (⊆→p (keep Δ⊆Γ))) E
h =
begin
  T.subst (T.exts (T.rename (⊆→p Δ⊆Γ) <$> T.id-subst)) <$> E
≡ ⟨ env-extensionality (<$>-fun (helper-5 Δ⊆Γ) E) ⟩
  _<$>_ {W = T.Lam} (T.rename (⊆→p (keep Δ⊆Γ))) E

```

```

N~N† : ∀ {Γ A} (N : S.Lam A Γ)
→ N ~ N†

```

```

N~N† (S.V x) with cc (S.V x)
N~N† (S.V x) | ∃[ Δ ] Δ⊆Γ ∧ N rewrite helper-2 x = ~V
N~N† (S.A M N) with cc M | cc N | inspect _† M | inspect _† N
N~N† (S.A M N) | ∃[ Δ1 ] Δ1⊆Γ ∧ M† | ∃[ Δ2 ] Δ2⊆Γ ∧ N†
| [ p ] | [ q ] with merge Δ1⊆Γ Δ2⊆Γ
N~N† (S.A M N) | ∃[ Δ1 ] Δ1⊆Γ ∧ M† | ∃[ Δ2 ] Δ2⊆Γ ∧ N†
| [ p ] | [ q ] | subListSum Γ1 Γ1⊆Γ Δ1⊆Γ1 Δ2⊆Γ1 well well1
rewrite helper-4 Δ1⊆Γ1 Γ1⊆Γ Δ1⊆Γ M† well
| helper-4 Δ2⊆Γ1 Γ1⊆Γ Δ2⊆Γ N† well1 | sym p | sym q

```

```

= ~A (N~N† M) (N~N† N)
N~N† (S.L N) with cc N | inspect _† N
N~N† (S.L N) | ∃[ Δ ] Δ⊆Γ ∧ N' | [ p ]
with adjust-context Δ⊆Γ
N~N† (S.L N) | ∃[ Δ ] Δ⊆Γ ∧ N' | [ p ]
| adjust Δ1 Δ1⊆Γ Δ⊆AΔ1 well = ~L g
where
h : T.subst (T.exts (T.rename (⊆→p Δ1⊆Γ) <$> T.id-subst))
    (T.rename (⊆→p Δ⊆AΔ1) N') ≡ T.rename (⊆→p Δ⊆Γ) N'
h =
begin
  T.subst (T.exts (T.rename (⊆→p Δ1⊆Γ) <$> T.id-subst))
    (T.rename (⊆→p Δ⊆AΔ1) N')
≡⟨ helper-5 Δ1⊆Γ (T.rename (⊆→p Δ⊆AΔ1) N') ⟩
  T.rename (⊆→p (keep Δ1⊆Γ)) (T.rename (⊆→p Δ⊆AΔ1) N')
≡⟨ rename◦rename (⊆→p Δ⊆AΔ1) (⊆→p (keep Δ1⊆Γ)) N' ⟩
  T.rename (select (⊆→p Δ⊆AΔ1) (⊆→p (keep Δ1⊆Γ))) N'
≡⟨ cong (λ e → T.rename e N')
    (env-extensionality (helper-3 Δ⊆AΔ1 (keep Δ1⊆Γ))) ⟩
  T.rename (⊆→p (⊆-trans Δ⊆AΔ1 (keep Δ1⊆Γ))) N'
≡⟨ cong (λ e → T.rename (⊆→p e) N') (sym well) ⟩
  T.rename (⊆→p Δ⊆Γ) N'
■
g : N ~ T.subst (T.exts (T.rename (⊆→p Δ1⊆Γ) <$> T.id-subst))
    (T.rename (⊆→p Δ⊆AΔ1) N')
g rewrite h | sym p = N~N† N

```

9.2 Compatibility relation is a bisimulation